

Machine-checked claims reserving in Lean 4: blueprint

Robert Sneiderman

August 30, 2026

Contents

1	Introduction	4
2	The chain-ladder estimators	5
2.1	Definitions	5
2.2	Deterministic identities	5
3	Mack 1999 equals Mack 1993	7
4	Bornhuetter-Ferguson on the chain-ladder pattern	8
5	The stochastic layer	9
5.1	Model	9
5.2	Mack's Theorem 2	9
5.3	Mack's Theorem 1	10
5.4	The variance assumptions and the estimation variance of $\hat{f}_k$	10
5.5	Unbiasedness of the variance estimator	11
5.6	Process variance and the exact MSEP (Theorem 3)	11
5.7	From independence across accident years to the $\mathcal{D}_k$ form	12
6	The variant catalogue: Mack versus conditional resampling	14
6.1	The total reserve: Mack's Corollary	14
7	Solvency II: the one-year view	16
7.1	The one-year claims development result	16
8	Röhr 2016: the error-propagation form	18
8.1	Definitions	18
8.2	The split, and the two Mack plug-ins	19
8.3	Catalogue row 3	19
8.4	The total reserve	20
8.5	The total reserve: the exact conditional MSEP	20
9	Mack 1999: general weights and exponent	22
9.1	Definitions	22
9.2	Deterministic identities	23
9.3	The recursion equals (*)	23
9.4	The 1993 case	24
9.5	Conditional unbiasedness of the weighted estimator	24

9.6	Weighted variance and estimator unbiasedness	25
10	The over-dispersed Poisson model	26
10.1	Definitions	26
10.2	The score equations are score equations	27
10.3	Chain ladder solves the score equations	27
10.4	Uniqueness: the marginal-sum solution is chain ladder	28
11	A finite independent-row witness	30
12	Merz-Wüthrich: the MSEP of the one-year claims development result	31
12.1	The exact one-year process variance	31
12.2	The estimator of Merz and Wüthrich	32
12.3	The bridge to the exact statement	33
12.4	Aggregation over accident years	33
13	The observed data: Mack's hypotheses discharged	34
13.1	From independence across accident years to the hypotheses	34
14	Mack's published numerical example	36
15	Mack's model diagnostics	37
15.1	Development-factor correlation	37
15.2	Calendar-year labels	37
16	Mack's lognormal approximation	39
17	Prior uncertainty in Bornhuetter-Ferguson	40
18	Bühlmann-Straub credibility	41
19	Munich chain ladder: the separate-chain-ladder gap	42
20	Munich residual correlations and adjusted factors	43
21	The stochastic ODP layer	45
21.1	Model and conditional moments	45
21.2	Prediction error	46
21.3	Log-link and corner constraints	46
21.4	Residual bootstrap definitions	47
21.5	Finite non-vacuity witness	47
22	Mack 1999: row assumptions and active weights	48
23	Calendar-year claims development results	50
24	Röhr: arbitrary-horizon approximations	52
24.1	Source formulas	52
24.2	Exact splits	53
24.3	Endpoints and bounds	53
25	R ChainLadder option semantics	54

26 Stochastic Bühlmann-Straub credibility	55
27 The stochastic Bornhuetter-Ferguson model	57
28 Panjer recursion	60
29 Non-vacuity: a nondegenerate model	61
30 Reference computation	62

Chapter 1

Introduction

This blueprint tracks formalized claims-reserving and actuarial risk results in Lean 4 on top of mathlib, with Mack’s distribution-free chain-ladder model [1] as the central strand. Each statement below carries the name of its Lean declaration; green nodes in the dependency graph are linked to checked declarations, while any open work is listed in the roadmap. Statements use source-aligned terms where applicable, and where a published formula is an estimator or an approximation rather than an identity, the text says so.

A run-off triangle with n accident years and n development years is a function $C : \mathbb{N} \rightarrow \mathbb{N} \rightarrow \mathbb{R}$; $C_{i,k}$ is the cumulative claims of accident year i at development year k , observed when $i+k \leq n-1$. The latest observed entry of accident year i is $C_{i,n-1-i}$.

Chapter 2

The chain-ladder estimators

2.1 Definitions

Definition 2.1 (Contributors). The accident years contributing to the development factor $k \rightarrow k+1$ are $i = 0, \dots, n-k-2$, those with $C_{i,k+1}$ observed.

Definition 2.2 (Column sums). $S_k = \sum_{i=0}^{n-k-2} C_{i,k}$ and $T_k = \sum_{i=0}^{n-k-2} C_{i,k+1}$.

Definition 2.3 (Development factor). $\hat{f}_k = T_k/S_k$. (In Lean, division by zero is zero.)

Definition 2.4 (Individual factor). $F_{i,k} = C_{i,k+1}/C_{i,k}$.

Definition 2.5 (Mack's variance estimator). $\hat{\sigma}_k^2 = \frac{1}{n-k-2} \sum_{i=0}^{n-k-2} C_{i,k} (F_{i,k} - \hat{f}_k)^2$. The value for $k = n-2$ (a single contributor) is an extrapolation convention in [1], not an estimator; no theorem here uses it.

Definition 2.6 (Projection). $\hat{C}_{i,k} = C_{i,n-1-i} \prod_{j=n-1-i}^{k-1} \hat{f}_j$ for $k \geq n-1-i$.

Definition 2.7 (Ultimate and reserve). $\hat{C}_{i,n-1}$ is the chain-ladder ultimate and $\hat{R}_i = \hat{C}_{i,n-1} - C_{i,n-1-i}$ the reserve.

Definition 2.8 (Mack's MSEP estimator). Mack (1993), Theorem 3:

$$\widehat{\text{mse}}(\hat{R}_i) = \hat{C}_{i,n-1}^2 \sum_{k=n-1-i}^{n-2} \frac{\hat{\sigma}_k^2}{\hat{f}_k^2} \left(\frac{1}{\hat{C}_{i,k}} + \frac{1}{S_k} \right).$$

2.2 Deterministic identities

Lemma 2.9 (Weighted-average form). If $C_{i,k} \neq 0$ for every contributor i , then $\hat{f}_k = \sum_i \frac{C_{i,k}}{S_k} F_{i,k}$.

Proof. Termwise: $\frac{C_{i,k}}{S_k} \cdot \frac{C_{i,k+1}}{C_{i,k}} = \frac{C_{i,k+1}}{S_k}$. □

Lemma 2.10. If $C_{i,k} \neq 0$ for every contributor, $T_k = \sum_i C_{i,k} F_{i,k}$.

Proof. Cancel the denominator termwise. □

Lemma 2.11 (Sum-of-squares decomposition). For any centre f , with $C_{i,k} \neq 0$ for every contributor, $\sum_i C_{i,k} (F_{i,k} - f)^2 = \sum_i C_{i,k} F_{i,k}^2 - 2fT_k + f^2S_k$.

Proof. Expand the square and use Lemma 2.10. □

Lemma 2.12 (Collapse at the chain-ladder centre). *If also $S_k \neq 0$, then $\sum_i C_{i,k}(F_{i,k} - \hat{f}_k)^2 = \sum_i C_{i,k}F_{i,k}^2 - S_k \hat{f}_k^2$. This is the identity behind the $n - k - 2$ degrees of freedom in $\hat{\sigma}_k^2$.*

Proof. Put $f = \hat{f}_k = T_k/S_k$ in Lemma 2.11. □

Lemma 2.13 (Projection step). *For $k \geq n - 1 - i$, $\hat{C}_{i,k+1} = \hat{C}_{i,k} \hat{f}_k$.*

Proof. Split off the top factor of the product. □

Lemma 2.14. $\hat{C}_{i,n-1-i} = C_{i,n-1-i}$.

Proof. Empty product. □

Lemma 2.15. *The oldest accident year is fully developed: $\hat{R}_0 = 0$.*

Proof. Empty product again. □

Chapter 3

Mack 1999 equals Mack 1993

Definition 3.1 (Mack's 1999 recursion). For $\alpha = 1$ and unit weights [2], starting from $\text{se}^2 = 0$ on the latest diagonal $d = n - 1 - i$ and stepping m times,

$$\text{se}_{m+1}^2 = \hat{C}_{i,d+m}^2 \left(\frac{\hat{\sigma}_{d+m}^2}{\hat{C}_{i,d+m}} + \frac{\hat{\sigma}_{d+m}^2}{S_{d+m}} \right) + \text{se}_m^2 \hat{f}_{d+m}^2.$$

Definition 3.2. The summand of the closed form: $\text{term}_{i,k} = \frac{\hat{\sigma}_k^2}{\hat{f}_k^2} \left(\frac{1}{\hat{C}_{i,k}} + \frac{1}{S_k} \right)$.

Lemma 3.3. $\widehat{\text{mse}}(\hat{R}_i) = \hat{C}_{i,n-1}^2 \sum_{k=n-1-i}^{n-2} \text{term}_{i,k}$.

Proof. By definition. □

Lemma 3.4 (Truncated closed form). If $\hat{f}_k \neq 0$ for $d \leq k < d+m$, then $\text{se}_m^2 = \hat{C}_{i,d+m}^2 \sum_{k=d}^{d+m-1} \text{term}_{i,k}$.

Proof. Induction on m using Lemma 2.13 and $\hat{f}_{d+m}^2 / \hat{f}_{d+m}^2 = 1$. □

Theorem 3.5 (Mack 1999 = Mack 1993). For $i \leq n - 1$ and $\hat{f}_k \neq 0$ along the row, i steps of the 1999 recursion return exactly Mack's 1993 estimator: $\text{se}_i^2 = \widehat{\text{mse}}(\hat{R}_i)$.

Proof. Lemma 3.4 with $m = i$, since $d + i = n - 1$. □

Chapter 4

Bornhuetter-Ferguson on the chain-ladder pattern

Definition 4.1 (Cumulative development factor, BF reserve and ultimate). $\text{cdf}_i = \prod_{k=d}^{n-2} \hat{f}_k$; for an a priori ultimate U , $R_i^{BF} = U(1 - 1/\text{cdf}_i)$ and $\hat{C}_{i,n-1}^{BF} = C_{i,d} + R_i^{BF}$.

Lemma 4.2 (Linearity). $R_i^{BF}(aU) = aR_i^{BF}(U)$ and $R_i^{BF}(U + V) = R_i^{BF}(U) + R_i^{BF}(V)$.

Proof. Ring identities. □

Theorem 4.3 (BF with the chain-ladder prior is chain ladder). *If $\text{cdf}_i \neq 0$ and $U = \hat{C}_{i,n-1}$, then $\hat{C}_{i,n-1}^{BF} = \hat{C}_{i,n-1}$ and $R_i^{BF} = \hat{R}_i$.*

Proof. $\hat{C}_{i,n-1} = C_{i,d} \text{cdf}_i$, so $U(1 - 1/\text{cdf}_i) = C_{i,d} \text{cdf}_i - C_{i,d}$. □

Lemma 4.4 (The BF reserve is a fraction of the prior). *If every development factor along the row is at least 1, then $\text{cdf}_i \geq 1$ and, for $U \geq 0$, $0 \leq R_i^{BF}(U) \leq U$.*

Proof. A product of numbers at least 1 is at least 1; then $0 \leq 1 - 1/\text{cdf}_i \leq 1$. □

Chapter 5

The stochastic layer

5.1 Model

Definition 5.1 (Random triangle). A random run-off triangle on a measurable space Ω is a family of random variables $C_{i,k} : \Omega \rightarrow \mathbb{R}$ together with a filtration $\mathcal{D}_k$ such that $C_{i,j}$ is $\mathcal{D}_k$ -measurable whenever $j \leq k$ ("everything observed up to development year k ").

Definition 5.2 (Assumption (M1)). $E[C_{i,k+1} \mid \mathcal{D}_k] = f_k C_{i,k}$ almost surely, for all i and k . This is Mack's first assumption in the form conditioned on $\mathcal{D}_k$; Mack states it conditioned on the row's own history and obtains this form from independence across accident years (M2). That passage is Theorem 5.30.

Definition 5.3 (Estimators as random variables). $S_k(\omega)$ and $\hat{f}_k(\omega)$ are the deterministic estimators applied to the triangle observed at ω .

Lemma 5.4. S_k is $\mathcal{D}_k$ -measurable.

Proof. A finite sum of $\mathcal{D}_k$ -measurable functions. □

5.2 Mack's Theorem 2

Theorem 5.5 (Conditional unbiasedness of the development factor). Assume (M1), that $S_k \neq 0$ almost surely, that each $C_{i,k+1}$ and $\hat{f}_k$ are integrable. Then $E[\hat{f}_k \mid \mathcal{D}_k] = f_k$ almost surely.

Proof. $\hat{f}_k = S_k^{-1} \sum_i C_{i,k+1}$. The factor S_k^{-1} is $\mathcal{D}_k$ -measurable and comes out of the conditional expectation; by (M1) the conditional expectation of the sum is $f_k \sum_i C_{i,k} = f_k S_k$; cancel S_k where it is nonzero. □

Lemma 5.6. $\hat{f}_j$ is $\mathcal{D}_k$ -measurable whenever $j + 1 \leq k$.

Proof. S_j is $\mathcal{D}_j \subseteq \mathcal{D}_k$ -measurable and each $C_{i,j+1}$ is $\mathcal{D}_k$ -measurable. □

Theorem 5.7 (Uncorrelated development factors, conditional form). Under (M1), on a finite measure space, for $j < k$ (with the integrability and $S \neq 0$ hypotheses of Theorem 5.5 for both columns and $\hat{f}_j \hat{f}_k$ integrable), $E[\hat{f}_j \hat{f}_k \mid \mathcal{D}_j] = f_j f_k$ almost surely.

Proof. Condition on $\mathcal{D}_k$ first: $\hat{f}_j$ is $\mathcal{D}_k$ -measurable and comes out, and $E[\hat{f}_k \mid \mathcal{D}_k] = f_k$ by Theorem 5.5. Then the tower property down to $\mathcal{D}_j$ and Theorem 5.5 again for column j . □

Theorem 5.8 (Mack 1993, Theorem 2). *On a probability space, under the same hypotheses, $E[\hat{f}_k] = f_k$ and $E[\hat{f}_j \hat{f}_k] = f_j f_k$ for $j < k$: the chain-ladder factors are unbiased and pairwise uncorrelated.*

Proof. Integrate the conditional statements. $\square$

5.3 Mack's Theorem 1

Definition 5.9 (Chain-ladder projection as a random variable). $\hat{C}_{i,d+m} = C_{i,d} \prod_{k=d}^{d+m-1} \hat{f}_k$ with $d = n - 1 - i$.

Lemma 5.10. $\hat{C}_{i,d+m}$ is $\mathcal{D}_{d+m}$ -measurable.

Proof. Induction on m ; each new factor $\hat{f}_{d+m}$ is $\mathcal{D}_{d+m+1}$ -measurable. $\square$

Lemma 5.11 (Iterating (M1)). *For $d \leq k$, $E[C_{i,k+1} \mid \mathcal{D}_d] = f_k E[C_{i,k} \mid \mathcal{D}_d]$, hence $E[C_{i,d+m} \mid \mathcal{D}_d] = C_{i,d} \prod_{k=d}^{d+m-1} f_k$.*

Proof. Tower property through $\mathcal{D}_k$, then (M1); induction on m . $\square$

Theorem 5.12 (Mack 1993, Theorem 1). *Under (M1), with the partial products integrable and the column sums a.s. nonzero, $E[\hat{C}_{i,d+m} \mid \mathcal{D}_d] = C_{i,d} \prod_{k=d}^{d+m-1} f_k$, and therefore $E[\hat{C}_{i,n-1} \mid \mathcal{D}_d] = E[C_{i,n-1} \mid \mathcal{D}_d]$: the chain-ladder ultimate is a conditionally unbiased estimator of the true ultimate claims.*

Proof. Condition on $\mathcal{D}_{d+m}$: the partial product comes out and the last factor has conditional expectation f_{d+m} (Theorem 5.5); tower down to $\mathcal{D}_d$ and induct. Compare with Lemma 5.11. $\square$

5.4 The variance assumptions and the estimation variance of $\hat{f}_k$

Definition 5.13 (Residuals, (M3) and (M2')). With $\varepsilon_{i,k} = C_{i,k+1} - f_k C_{i,k}$: (M3) $E[\varepsilon_{i,k}^2 \mid \mathcal{D}_k] = \sigma_k^2 C_{i,k}$, and (M2') $E[\varepsilon_{i,k} \varepsilon_{j,k} \mid \mathcal{D}_k] = 0$ for $i \neq j$, the conditional uncorrelatedness across accident years that independence of accident years supplies.

Lemma 5.14. *On $\{S_k \neq 0\}$, $\hat{f}_k - f_k = S_k^{-1} \sum_i \varepsilon_{i,k}$ and $(\hat{f}_k - f_k)^2 = S_k^{-2} \sum_i \sum_j \varepsilon_{i,k} \varepsilon_{j,k}$.*

Proof. $T_k - f_k S_k = \sum_i \varepsilon_{i,k}$; expand the square of the sum. $\square$

Theorem 5.15 (Estimation variance of the development factor). *Under (M3) and (M2'), with the residual products integrable, on $\{S_k \neq 0\}$, $E[(\hat{f}_k - f_k)^2 \mid \mathcal{D}_k] = \sigma_k^2 / S_k$ (Mack 1993, proof of Theorem 3; Mack 1999, s.e. $(\hat{f}_k)^2 = \hat{\sigma}_k^2 / S_k$).*

Proof. S_k^{-2} is $\mathcal{D}_k$ -measurable and comes out; conditional expectation is linear over the double sum; cross terms vanish by (M2'); diagonal terms give $\sigma_k^2 \sum_i C_{i,k} = \sigma_k^2 S_k$. $\square$

5.5 Unbiasedness of the variance estimator

Lemma 5.16 (Sum of squares around $\hat{f}_k$ via residuals). *With all contributing $C_{i,k} \neq 0$ and $S_k \neq 0$, for any f , $\sum_i C_{i,k}(F_{i,k} - \hat{f}_k)^2 = \sum_i \varepsilon_i^2 / C_{i,k} - S_k(\hat{f}_k - f)^2$ where $\varepsilon_i = C_{i,k+1} - fC_{i,k}$.*

Proof. Two instances of Lemma 2.11 (centres $\hat{f}_k$ and f) and $T_k = \hat{f}_k S_k$. $\square$

Definition 5.17. $\hat{\sigma}_k^2$ and the weighted sum of squares $W_k = \sum_i C_{i,k}(F_{i,k} - \hat{f}_k)^2$ as random variables; $\hat{\sigma}_k^2 = W_k / (n - k - 2)$.

Theorem 5.18 (Unbiasedness of the variance estimator). *Under (M3) and (M2'), with all contributing $C_{i,k}$ and S_k a.s. nonzero and the stated integrability, $E[W_k | \mathcal{D}_k] = (n - k - 2) \sigma_k^2$ for $k + 2 \leq n$, and hence $E[\hat{\sigma}_k^2 | \mathcal{D}_k] = \sigma_k^2$ for $k + 3 \leq n$.*

Proof. By Lemma 5.16, $W_k = \sum_i \varepsilon_i^2 / C_{i,k} - S_k(\hat{f}_k - f_k)^2$. Each $C_{i,k}^{-1}$ is $\mathcal{D}_k$ -measurable, so (M3) gives σ_k^2 per term, $n - k - 1$ terms in all; S_k is $\mathcal{D}_k$ -measurable and Theorem 5.15 gives σ_k^2 for the subtracted term. $\square$

5.6 Process variance and the exact MSEP (Theorem 3)

Lemma 5.19 (Conditional second moment). *Under (M1) and (M3), $E[\varepsilon_{i,k} | \mathcal{D}_k] = 0$, $E[C_{i,k+1}^2 | \mathcal{D}_k] = f_k^2 C_{i,k}^2 + \sigma_k^2 C_{i,k}$, and for $d \leq k$, $E[C_{i,k+1}^2 | \mathcal{D}_d] = f_k^2 E[C_{i,k}^2 | \mathcal{D}_d] + \sigma_k^2 E[C_{i,k} | \mathcal{D}_d]$.*

Proof. Expand $C_{i,k+1} = f_k C_{i,k} + \varepsilon_{i,k}$; the cross term vanishes; tower property. $\square$

Definition 5.20 (Process variance). $V_0 = 0$, $V_{m+1} = f_{d+m}^2 V_m + \sigma_{d+m}^2 c \prod_{l < m} f_{d+l}$; in closed form $V_m = \sum_{j < m} \left(\prod_{l=j+1}^{m-1} f_{d+l} \right)^2 \sigma_{d+j}^2 c \prod_{l < j} f_{d+l}$.

Theorem 5.21 (Process variance along a row). *For $i < n$, under (M1), (M3), and the stated claim, claim-square, residual-square, and claim-residual integrability hypotheses, $E[C_{i,d+m}^2 | \mathcal{D}_d] - E[C_{i,d+m} | \mathcal{D}_d]^2 = V_m$ with $c = C_{i,d}$.*

Proof. Induction on m with Lemma 5.19 and Lemma 5.11. $\square$

Lemma 5.22 (Conditional MSEP decomposition). *For a sub-sigma-algebra $\mathcal{D}$, a $\mathcal{D}$ -measurable predictor P of Y , and the stated target-square, predictor-square, and product integrability hypotheses, $E[(P - Y)^2 | \mathcal{D}] = (E[Y^2 | \mathcal{D}] - E[Y | \mathcal{D}]^2) + (P - E[Y | \mathcal{D}])^2$.*

Proof. Expand the square; P comes out of the conditional expectation. $\square$

Theorem 5.23 (Mack's Theorem 3, exact form). *Let $\mathcal{D}$ be the observed data and $i < n$. Assume $\hat{C}_{i,d+m}$ is $\mathcal{D}$ -measurable and that $\mathcal{D}$ adds nothing to $\mathcal{D}_d$ about row i 's future (the conditional mean and second moment of $C_{i,d+m}$ given $\mathcal{D}$ agree with those given $\mathcal{D}_d$; this is what independence across accident years supplies). Then under (M1), (M3), and the stated claim, residual, predictor, square, and product integrability hypotheses,*

$$E[(\hat{C}_{i,d+m} - C_{i,d+m})^2 | \mathcal{D}] = V_m + (\hat{C}_{i,d+m} - C_{i,d} \prod_{l < m} f_{d+l})^2.$$

Proof. Lemma 5.22 with $P = \hat{C}_{i,d+m}$, then Theorem 5.21 and Lemma 5.11. $\square$

Definition 5.24 (Mack’s estimators of the two terms). $\hat{V}$ is V_m with $\hat{f}, \hat{\sigma}^2$ plugged in. The estimation-error term $\hat{C}_{i,n-1}^2 \sum_k \hat{\sigma}_k^2 / (\hat{f}_k^2 S_k)$ is obtained from $(\hat{C} - M)^2$ by Mack’s conditional-resampling step: replace $(\hat{f}_k - f_k)^2$ by σ_k^2 / S_k (Theorem 5.15), drop the cross terms (Theorem 5.7), and plug in $\hat{f}, \hat{\sigma}^2$. This is a definition, not a theorem: it is the approximation in Mack [1], and the exact conditional MSEP above is the object it approximates. Buchwalder, Bühlmann, Merz and Wüthrich [15] estimate the same object differently.

5.7 From independence across accident years to the $\mathcal{D}_k$ form

Mack states (M1) and (M3) conditioned on one accident year’s own history, and assumes separately that the accident years are independent (M2). Every theorem above conditions on $\mathcal{D}_k$ instead. This section proves the passage.

Definition 5.25 (Row σ -algebras). $\sigma(C_{i,0}, \dots, C_{i,k})$, the history of accident year i up to development year k ; $\sigma(C_{i,j} : j)$, the whole of accident year i ; and the joins of these over the other accident years $i' \neq i$.

Definition 5.26 ((M2) and the row-conditioned assumptions). (M2) is **RowsIndep**: the σ -algebras of the individual accident years are independent. **RowsGenerated** says that $\mathcal{D}_k$ is what its name means, the join of the rows’ histories up to development year k . (M1row) and (M3row) are Mack’s assumptions conditioned on $\sigma(C_{i,0}, \dots, C_{i,k})$ rather than on $\mathcal{D}_k$.

Lemma 5.27 (Rectangles). *If m'_1 and m_2 are independent, F is m'_1 -measurable and integrable, $a \in m'_1$ and $b \in m_2$, then $\int_{a \cap b} F = \mu(b) \int_a F$.*

Proof. Apply `mathlib’s condExp_indep_eq` to the indicator $1_a F$, which is m'_1 -measurable, and integrate the resulting constant over b . $\square$

Theorem 5.28 (Conditioning on an independent enlargement changes nothing). *Let $m_1 \leq m'_1$, let f be m'_1 -measurable and integrable, and let m_2 be independent of m'_1 . Then $E[f \mid m_1 \vee m_2] = E[f \mid m_1]$ almost surely.*

Proof. The sets $a \cap b$ with $a \in m_1$, $b \in m_2$ form a π -system generating $m_1 \vee m_2$. On such a set both $\int f$ and $\int E[f \mid m_1]$ factor by Lemma 5.27 into $\mu(b)$ times the integral over a , and those agree by the defining property of $E[f \mid m_1]$. The class of sets on which the two set integrals agree is closed under complements and countable disjoint unions, so it is all of $m_1 \vee m_2$; conclude by the uniqueness of the conditional expectation. `Mathlib` has the case $m_1 = \perp$ only. $\square$

Lemma 5.29 (One row against the others). *Under (M2), accident year i is independent of all the other accident years taken together, and so of their observed histories. Under **RowsGenerated**, $\mathcal{D}_k$ is the join of row i ’s history up to k with the other rows’ histories up to k .*

Proof. `indep_iSup_of_disjoint` for $\{i\}$ against its complement, then monotonicity of independence; the splitting of $\mathcal{D}_k$ is `iSup_split_single`. $\square$

Theorem 5.30 ((M1) and (M3) in the $\mathcal{D}_k$ form). *Assume (M2), that $\mathcal{D}_k$ is generated by the rows’ histories, and the integrability of the entries. Then Mack’s row-conditioned (M1row) implies (M1), and (M3row) implies (M3).*

Proof. Split $\mathcal{D}_k$ by Lemma 5.29, apply Theorem 5.28 with m_1 the row’s history up to k , m'_1 the whole row and m_2 the other rows, and finish with the row-conditioned assumption. $\square$

Theorem 5.31 ((M2') is a consequence, not an assumption). *Under (M2), RowsGenerated and (M1row), the residuals of two different accident years are conditionally uncorrelated given $\mathcal{D}_k$: assumption (M2'), which the variance theorems take as a hypothesis, follows from independence across accident years.*

Proof. Condition on $\mathcal{D}_k$ together with the whole of accident year j : row j 's residual comes out of the conditional expectation, and what is left is row i 's residual conditioned on a σ -algebra that adds nothing about row i , so it equals $E[\varepsilon_{i,k} \mid \sigma(C_{i,0}, \dots, C_{i,k})] = 0$ by Theorem 5.28 and (M1row). The tower property brings the identity back to $\mathcal{D}_k$. $\square$

Chapter 6

The variant catalogue: Mack versus conditional resampling

Write $a_k = \hat{\sigma}_k^2 / (\hat{f}_k^2 S_k)$ for the relative estimation variance of development factor k along the row. Mack's estimation-error term is $\hat{C}_{i,n-1}^2 \sum_k a_k$; the conditional-resampling term of Buchwalder, Bühlmann, Merz and Wüthrich [15], which Murphy [16] also reaches, is $\hat{C}_{i,n-1}^2 (\prod_k (1 + a_k) - 1)$.

Definition 6.1. a_k and the conditional-resampling estimator $\hat{C}_{i,n-1}^2 (\prod_k (1 + a_k) - 1)$.

Lemma 6.2 (Product minus sum). *For $a_k \geq 0$: $1 + \sum a_k \leq \prod (1 + a_k) \leq \exp(\sum a_k)$; for two terms $(1 + a)(1 + b) - 1 - (a + b) = ab$; and $\prod (1 + a_k) - 1 - \sum a_k = 0$ when at most one a_k is nonzero.*

Proof. Induction over the index set for the Weierstrass inequality; $1 + x \leq e^x$ termwise for the exponential bound. $\square$

Theorem 6.3 (Catalogue row: Mack is the first-order part of BMW). *With $a_k \geq 0$ along the row: (i) Mack $\leq$ BMW; (ii) the difference is exactly $\hat{C}_{i,n-1}^2 (\prod_k (1 + a_k) - 1 - \sum_k a_k)$; (iii) the two coincide for the second-latest accident year, which has a single development factor; (iv) the difference is at most $\hat{C}_{i,n-1}^2 (e^{\sum a_k} - 1 - \sum a_k)$, second order in the relative estimation variances.*

Proof. Lemma 6.2 scaled by $\hat{C}_{i,n-1}^2$. $\square$

Theorem 6.4 (Strict counterexample). *There is a run-off triangle (four accident years, four development years, with unequal individual factors in the first two columns) and an accident year on which Mack's estimation-error term is strictly smaller than the conditional-resampling term: the two are different estimators, not two presentations of one. On the exhibited triangle $a_0 = 1/297$, $a_1 = 1/2809$, and the remainder is $\hat{C}^2 a_0 a_1 > 0$. Every number is checked by `norm_num` from the definitions.*

Proof. Evaluate the definitions on the triangle; apply Theorem 6.3(ii). $\square$

6.1 The total reserve: Mack's Corollary

Definition 6.5 (Total-reserve estimators). Mack's Corollary: $\widehat{\text{mse}}(\hat{R}) = \sum_i \widehat{\text{mse}}_i + \sum_i \hat{C}_i (\sum_{j>i} \hat{C}_j) 2 \sum_{k \in \text{row } i} a_k$, the cross terms running over the older year's remaining factors. The conditional-resampling aggregate replaces $2 \sum_k a_k$ by $2(\prod_k (1 + a_k) - 1)$ [15, Section 4.3].

Theorem 6.6 (Aggregated catalogue row). *The aggregated estimation-error terms differ by $\sum_i (\hat{C}_i^2 + 2\hat{C}_i \sum_{j>i} \hat{C}_j) (\prod_{k \in \text{row } i} (1 + a_k) - 1 - \sum_{k \in \text{row } i} a_k)$, which is nonnegative when the ultimates and the a_k are: Mack's total is at most the resampling total, by the same row-wise second-order remainder as in the single-year case.*

Proof. Termwise algebra and Lemma 6.2. □

Chapter 7

Solvency II: the one-year view

7.1 The one-year claims development result

Solvency II prices the uncertainty of the next twelve months, not of the run-off to ultimate. Merz and Wüthrich [7] call the twelve-month move of the predicted ultimate the claims development result and separate the true one from the observable one. The filtration here is the one carried by `RandomTriangle`: $\mathcal{D}_k$ is everything observed up to development year k , for every accident year, the filtration in which (M1) is assumed in this development. Merz and Wüthrich index information by calendar year instead; the martingale property below is the same tower property in either filtration, and the closed form is stated in the filtration (M1) is assumed in.

Definition 7.1 (True one-year claims development result). For accident year i and development year k ,

$$\text{CDR}_i(k) = E[C_{i,n-1} \mid \mathcal{D}_k] - E[C_{i,n-1} \mid \mathcal{D}_{k+1}],$$

the change in the predicted ultimate claim when the information grows by one year.

Theorem 7.2 (Merz-Wüthrich 2008: the true CDR has conditional mean zero). *On a finite measure space, $E[\text{CDR}_i(k) \mid \mathcal{D}_k] = 0$ almost surely, for every accident year i and every k : the reserve set at time k is unbiased for the reserve that will be set at time $k + 1$ together with the payments made in between.*

Proof. Linearity of the conditional expectation, then the tower property twice: the inner $\mathcal{D}_k$ -conditional expectation is idempotent, and $E[E[C_{i,n-1} \mid \mathcal{D}_{k+1}] \mid \mathcal{D}_k] = E[C_{i,n-1} \mid \mathcal{D}_k]$ because $\mathcal{D}_k \subseteq \mathcal{D}_{k+1}$. The two terms cancel. The only hypotheses are that $\mathcal{D}$ is a filtration inside the ambient σ -algebra and that the measure is finite; no integrability assumption on $C_{i,n-1}$ is needed for the identity as stated, since both conditional expectations are 0 off the integrable case. $\square$

Lemma 7.3 (Iterating (M1) from an arbitrary development year). *Under (M1), $E[C_{i,d+m} \mid \mathcal{D}_d] = C_{i,d} \prod_{j=d}^{d+m-1} f_j$ for every starting year d , hence $E[C_{i,n-1} \mid \mathcal{D}_k] = C_{i,k} \prod_{j=k}^{n-2} f_j$ for $k \leq n - 1$.*

Proof. The induction of Lemma 5.11 with the starting year free rather than pinned to the latest observed year $n - 1 - i$ of the row; the base case is that $C_{i,d}$ is $\mathcal{D}_d$ -measurable. $\square$

Theorem 7.4 (Merz-Wüthrich 2008: the true CDR as a one-step residual). *Under (M1), for $k < n - 1$ and integrable claims,*

$$\text{CDR}_i(k) = \left(\prod_{j=k+1}^{n-2} f_j \right) (f_k C_{i,k} - C_{i,k+1})$$

almost surely: everything the next year of experience does to the predicted ultimate is carried by the single deviation of the new observation from its (M1) prediction.

Proof. Both predicted ultimates are given by Lemma 7.3; splitting $\prod_{j=k}^{n-2} f_j = f_k \prod_{j=k+1}^{n-2} f_j$ at the bottom index leaves the common factor $\prod_{j>k} f_j$ times $f_k C_{i,k} - C_{i,k+1}$. $\square$

Definition 7.5 (Observable one-year claims development result). The difference of the two chain-ladder predictions an actuary computes: the ultimate estimated from the triangle with n diagonals minus the ultimate estimated a year later, with every development factor re-estimated, from the triangle with $n + 1$ diagonals, $\widehat{\text{CDR}}_i = \hat{C}_{i,n-1}^{(n)} - \hat{C}_{i,n-1}^{(n+1)}$, together with its value along a random triangle outcome by outcome.

Lemma 7.6 (The displayed form). $\widehat{\text{CDR}}_i = \hat{R}_i^{(n)} - ((C_{i,n-i} - C_{i,n-1-i}) + \hat{R}_i^{(n+1)})$: *opening reserve minus the payments of the year and the closing reserve, the form Merz and Wüthrich display. Algebra on the definitions.*

Proof. Expand both reserves; the two claim entries cancel. $\square$

Nothing is proved here about the distribution of the observable CDR. The conditional mean squared error of prediction of $\widehat{\text{CDR}}_i$, which is what a Solvency II one-year reserve risk figure reports, rests in Merz and Wüthrich on an approximation of the same kind as Mack's: the estimated development factors are treated as resampled ones and terms of second order in the residuals are dropped. That step is not formalized, so the true-CDR statements above are exact and the observable-CDR statements of the paper are not yet available here.

Chapter 8

Röhr 2016: the error-propagation form

Röhr, *Chain ladder and error propagation* [18], derives the chain-ladder prediction error from the error-propagation formula. The chain-ladder ultimate $\hat{C}_{i,n-1} = C_{i,d} \prod_{k=d}^{n-2} \hat{f}_k$, $d = n - 1 - i$, is a product, so a first-order expansion in the development factors turns the relative squared error of the product into the *sum* of the relative squared errors of the factors. The paper’s classical-case display is

$$\frac{\widehat{\text{mse}}(\hat{R}_i)}{\hat{C}_{i,n-1}^2} = \sum_k \hat{u}_k^2,$$

where $\hat{u}_k^2$ is the relative uncertainty attached to development step k , and the derivation splits it into a process part and a parameter part,

$$\hat{u}_k^2 = \frac{\hat{\sigma}_k^2}{\hat{f}_k^2 \hat{C}_{i,k}} + \frac{\hat{\sigma}_k^2}{\hat{f}_k^2 S_k},$$

the first from the conditional variance $\sigma_k^2 C_{i,k}$ of the step itself, the second from the estimation variance σ_k^2 / S_k of $\hat{f}_k$.

Source note. The full text of Röhr [18] is paywalled and was not available for this formalization. What is formalized is the classical-case (ultimate run-off, single accident year) display quoted above, taken from the paper’s own abstract — “in the classical case treated by Mack (1993) the mean squared prediction error divided by the squared estimated ultimate loss can be written as $\sum_j \hat{u}_j^2$, where $\hat{u}_j$ measures the (relative) uncertainty around the j -th development factor and the proportion of the estimated ultimate loss that it affects”, with “a split into process error and parameter error” — together with the two parts in the form above. A primary-author CAE 2014 presentation separately states the aggregate-influence formulas for arbitrary horizons; the chapter on arbitrary-horizon approximations records slides 13, 17, and 20 without attributing further formulas to the full paper. Röhr’s own aggregation over accident years remains unchecked.

8.1 Definitions

Definition 8.1 (Röhr’s relative uncertainties). $\hat{u}_{k,\text{proc}}^2 = \hat{\sigma}_k^2 / (\hat{f}_k^2 \hat{C}_{i,k})$, $\hat{u}_{k,\text{par}}^2 = \hat{\sigma}_k^2 / (\hat{f}_k^2 S_k)$, and $\hat{u}_k^2 = \hat{u}_{k,\text{proc}}^2 + \hat{u}_{k,\text{par}}^2$.

Definition 8.2 (Röhr's linearized prediction error). The process term $\hat{C}_{i,n-1}^2 \sum_{k=d}^{n-2} \hat{u}_{k,\text{proc}}^2$, the parameter term $\hat{C}_{i,n-1}^2 \sum_{k=d}^{n-2} \hat{u}_{k,\text{par}}^2$, and their sum $\text{rohrMsep}_i = \hat{C}_{i,n-1}^2 \sum_{k=d}^{n-2} \hat{u}_k^2$.

8.2 The split, and the two Mack plug-ins

Lemma 8.3 (Process plus parameter). rohrMsep_i is the process term plus the parameter term.

Proof. Termwise, by construction. $\square$

Lemma 8.4 (Parameter term is Mack's estimation term). $\hat{u}_{k,\text{par}}^2 = a_k$, so Röhr's parameter term is Mack's estimation-error term $\hat{C}_{i,n-1}^2 \sum_k a_k$ exactly.

Proof. $\hat{\sigma}_k^2 / (\hat{f}_k^2 S_k) = (\hat{\sigma}_k^2 / \hat{f}_k^2) / S_k$. $\square$

Lemma 8.5 (Splitting the projection). For $n-1-i \leq k \leq m$, $\hat{C}_{i,m} = \hat{C}_{i,k} \prod_{l=k}^{m-1} \hat{f}_l$.

Proof. Split the product of the definition at k . $\square$

Theorem 8.6 (Process term is Mack's process variance). Let $i \leq n-1$ and suppose $\hat{f}_k \neq 0$ and $\hat{C}_{i,k} \neq 0$ for every $k = d, \dots, n-2$. Then Röhr's process term equals procVar , Mack's process variance recursion $V_{m+1} = \hat{f}_{d+m}^2 V_m + \hat{\sigma}_{d+m}^2 \hat{C}_{i,d+m}$, $V_0 = 0$, run for i steps with $\hat{f}, \hat{\sigma}^2$ plugged in.

Proof. Write $\hat{C}_{i,n-1} = \hat{C}_{i,k} \hat{f}_k \prod_{l=k+1}^{n-2} \hat{f}_l$ by Lemma 8.5. Then $\hat{C}_{i,n-1}^2 \hat{\sigma}_k^2 / (\hat{f}_k^2 \hat{C}_{i,k}) = (\prod_{l=k+1}^{n-2} \hat{f}_l)^2 \hat{\sigma}_k^2 \hat{C}_{i,k}$, which is the k -th summand of the closed form of procVar ; reindex $k = d+j$, $j = 0, \dots, i-1$. $\square$

8.3 Catalogue row 3

Theorem 8.7 (Röhr = Mack 1993). $\text{rohrMsep}_i = \widehat{\text{mse}}(\hat{R}_i)$, with no side condition: the linearized error-propagation form and Mack's 1993 closed form are equal term by term. This catalogue row therefore has zero correction term, in contrast with Theorem 6.3. Dividing by $\hat{C}_{i,n-1}^2 \neq 0$ recovers Röhr's display $\widehat{\text{mse}}(\hat{R}_i) / \hat{C}_{i,n-1}^2 = \sum_k \hat{u}_k^2$.

Proof. $\hat{\sigma}_k^2 / (\hat{f}_k^2 \hat{C}_{i,k}) + \hat{\sigma}_k^2 / (\hat{f}_k^2 S_k) = (\hat{\sigma}_k^2 / \hat{f}_k^2) (1 / \hat{C}_{i,k} + 1 / S_k)$ termwise. $\square$

Corollary 8.8 (Mack's closed form is process plus estimation). Under the hypotheses of Theorem 8.6, $\widehat{\text{mse}}(\hat{R}_i) = \text{mackProcess}_i + \text{mackEstimation}_i$: reading Mack's formula through Röhr's split identifies its two halves with the two plug-in estimators of the stochastic layer.

Proof. Chain Theorem 8.7, Lemma 8.3, Theorem 8.6 and Lemma 8.4. $\square$

Theorem 8.9 (What the linearization costs). Röhr's parameter term is the first-order part of the conditional-resampling term: their difference is exactly $\hat{C}_{i,n-1}^2 (\prod_k (1+a_k) - 1 - \sum_k a_k)$, and with $a_k \geq 0$ along the row Röhr's term is the smaller one. Linearizing in the estimated factors is precisely the step that replaces the product by the sum.

Proof. Lemma 8.4 and Theorem 6.3. $\square$

8.4 The total reserve

Definition 8.10 (Aggregate). Röhr's single-year terms carried into Mack's Corollary, with Mack's cross terms: $\sum_i \text{rohrMsep}_i + \sum_i \text{mackCross}_i$. The cross terms are Mack's; the paper's own aggregation over accident years could not be checked against the source, so this is recorded as a definition rather than attributed.

Theorem 8.11. *That aggregate is Mack's total-reserve estimator $\widehat{\text{mse}}(\hat{R})$.*

Proof. Theorem 8.7 row by row. □

8.5 The total reserve: the exact conditional MSEP

Mack's Corollary estimates the mean squared error of prediction of the total reserve $\hat{R} = \sum_i \hat{R}_i$. The estimator is Definition 6.5; the object it estimates is the conditional MSEP of the sum, and this section proves what that object is.

Write $\hat{C}_i = \hat{C}_{i,n-1}$ for the chain-ladder ultimate of accident year i , $C_i = C_{i,n-1}$ for the true ultimate, $\mathcal{D}$ for the σ -algebra of the observed data (the one Theorem 5.23 conditions on) and $M_i = E[C_i \mid \mathcal{D}]$.

Definition 8.12 (Cross-term condition). For a finite set s of accident years and targets Y_i ,

$$E[(Y_i - E[Y_i \mid \mathcal{D}])(Y_j - E[Y_j \mid \mathcal{D}]) \mid \mathcal{D}] = 0 \quad \text{for } i \neq j \text{ in } s.$$

This is what independence across accident years supplies, in the same way that the $\mathcal{D}_k$ -conditioned form of (M1) (Definition 5.2) is what it supplies for a single row. Like (M1), (M3) and the “no extra information” hypotheses of Theorem 5.23, it is carried as a hypothesis here; the passage from independence of the rows to it is Theorem 13.6.

Lemma 8.13 (Conditional variance as a centred second moment). *For $Y \in L^2$, $E[(E[Y \mid \mathcal{D}] - Y)^2 \mid \mathcal{D}] = E[Y^2 \mid \mathcal{D}] - E[Y \mid \mathcal{D}]^2$.*

Proof. Lemma 5.22 with the predictor taken to be $E[Y \mid \mathcal{D}]$ itself, so that the estimation-error term is zero. Square integrability supplies every product integrability side condition by Hölder. □

Lemma 8.14 (Conditional variances add). *If $E[Z_i Z_j \mid \mathcal{D}] = 0$ for $i \neq j$ in s and each $Z_i \in L^2$, then $E[(\sum_i Z_i)^2 \mid \mathcal{D}] = \sum_i E[Z_i^2 \mid \mathcal{D}]$.*

Proof. Expand the square into the double sum $\sum_{i,j} Z_i Z_j$, take conditional expectations term by term, and drop the off-diagonal terms. □

Theorem 8.15 (Conditional MSEP of a sum). *For $\mathcal{D}$ -measurable predictors P_i of square-integrable targets Y_i whose centred versions satisfy Definition 8.12,*

$$E\left[\left(\sum_i P_i - \sum_i Y_i\right)^2 \mid \mathcal{D}\right] = \sum_i (E[Y_i^2 \mid \mathcal{D}] - E[Y_i \mid \mathcal{D}]^2) + \left(\sum_i (P_i - E[Y_i \mid \mathcal{D}])\right)^2.$$

The conditional variances add; the estimation errors add before they are squared.

Proof. Lemma 5.22 applied to the totals, then Lemma 8.13 to rewrite the conditional variance of the total as a centred second moment, then Lemma 8.14 to split it, then Lemma 8.13 again on each accident year. □

Lemma 8.16 (Theorem 1 and the process variance at the ultimate). *Under (M1), (M3) and the hypotheses of Theorem 5.23 that $\mathcal{D}$ carries no information about row i 's future beyond $\mathcal{D}_{n-1-i}$,*

$$M_i = C_{i,n-1-i} \prod_{l=n-1-i}^{n-2} f_l, \quad E[C_i^2 \mid \mathcal{D}] - M_i^2 = \text{procVar}(C_{i,n-1-i}, f, \sigma^2, n-1-i)_i.$$

Proof. Specialize Lemma 5.11 and Theorem 5.21 to $m = i$, where $(n-1-i) + i = n-1$, and transport them from $\mathcal{D}_{n-1-i}$ to $\mathcal{D}$. $\square$

Theorem 8.17 (Mack's Corollary, exact form). *Under (M1), (M3), square integrability of the claims and of the chain-ladder predictions, the "no extra information" hypotheses of Theorem 5.23 for each accident year, and Definition 8.12 for the true ultimates,*

$$E\left[\left(\sum_i \hat{C}_i - \sum_i C_i\right)^2 \mid \mathcal{D}\right] = \sum_i \text{procVar}_i + \left(\sum_i (\hat{C}_i - M_i)\right)^2.$$

The process variances add; the estimation errors add before being squared, and expanding that square is the exact origin of the cross terms of Definition 6.5.

Proof. Theorem 8.15 with $P_i = \hat{C}_i$ and $Y_i = C_i$, then Lemma 8.16 term by term. $\square$

Corollary 8.18 (The total MSEP dominates the process variance). $\sum_i \text{procVar}_i \leq E[(\sum_i \hat{C}_i - \sum_i C_i)^2 \mid \mathcal{D}]$ almost everywhere, whatever the chain-ladder predictions happen to be.

Proof. The estimation-error term of Theorem 8.17 is a square. $\square$

Theorem 8.19 (The plug-in has the same two-part shape). *On a row whose development factors $\hat{f}_k$ and projections $\hat{C}_{i,k}$ do not vanish, the $1/\hat{C}_{i,k}$ half of Mack's formula is exactly procVar with $\hat{f}, \hat{\sigma}^2$ substituted and the $1/S_k$ half is exactly the plug-in estimation error: $\widehat{\text{mse}}_i = \text{mackProcess}_i + \text{mackEstimation}_i$, hence $\widehat{\text{mse}}(\hat{R}) = \sum_i \text{mackProcess}_i + \text{mackTotalEstimation}$. This is the plug-in counterpart of Theorem 8.17: process variances added, and everything else, single-year terms and cross terms alike, estimating the squared summed estimation error.*

Proof. Reindex the row sum over $\text{Ico}(n-1-i, n-1)$ by $j < i$, split the product $\prod_{l=n-1-i}^{n-2} \hat{f}_l$ at the factor $\hat{f}_{n-1-i+j}$, and cancel; then sum over accident years. $\square$

Theorem 8.20 (The cross-term condition is not vacuous). *On the nondegenerate model of Theorem 29.1 the three true ultimates are genuinely random and their centred versions are conditionally uncorrelated given the trivial σ -algebra: Definition 8.12 holds for $s = \{0, 1, 2\}$ without holding for want of randomness.*

Proof. The unconditional mean of each ultimate is cf^2 ; the centred ultimates are proportional to the shocks ξ_i , which are orthogonal. $\square$

Chapter 9

Mack 1999: general weights and exponent

[2] generalizes the development factor of [1] to

$$\hat{f}_k = \frac{\sum_i w_{ik} C_{ik}^\alpha F_{ik}}{\sum_i w_{ik} C_{ik}^\alpha}, \quad F_{ik} = \frac{C_{i,k+1}}{C_{ik}},$$

with weights $w_{ik} \in [0, 1]$ and an exponent α . The three cases the paper names are $\alpha = 1$ with unit weights (the 1993 estimator), $\alpha = 0$ (the simple average of the individual factors) and $\alpha = 2$ (the ordinary regression of $C_{i,k+1}$ on C_{ik} through the origin). The variance assumption becomes $\text{Var}(F_{ik} \mid \cdot) = \sigma_k^2 / (w_{ik} C_{ik}^\alpha)$ and the variance estimator

$$\hat{\sigma}_k^2 = \frac{1}{n - k - 2} \sum_i w_{ik} C_{ik}^\alpha (F_{ik} - \hat{f}_k)^2.$$

Two displays of that paper are formalized here. The first is its equation (*),

$$\text{s.e.}(\hat{C}_{in})^2 = \hat{C}_{in}^2 \sum_k \frac{\hat{\sigma}_k^2}{\hat{f}_k^2} \left(\frac{1}{w_{ik} \hat{C}_{ik}^\alpha} + \frac{1}{\sum_j w_{jk} C_{jk}^\alpha} \right),$$

the second is the recursion displayed immediately below it,

$$\text{s.e.}(\hat{C}_{i,k+1})^2 = \hat{C}_{ik}^2 (\text{s.e.}(F_{ik})^2 + \text{s.e.}(\hat{f}_k)^2) + \text{s.e.}(\hat{C}_{ik})^2 \hat{f}_k^2, \quad \text{s.e.}(F_{ik})^2 = \frac{\hat{\sigma}_k^2}{w_{ik} \hat{C}_{ik}^\alpha}, \quad \text{s.e.}(\hat{f}_k)^2 = \frac{\hat{\sigma}_k^2}{\sum_j w_{jk} C_{jk}^\alpha}.$$

Theorem 3.5 proves the two agree for $\alpha = 1$ and unit weights; the theorem below carries the weights and the exponent through, and the unit case recovers the earlier statement. Mack's restriction $w_{ik} \in [0, 1]$ and $\alpha \in \{0, 1, 2\}$ is a modelling choice and no identity here uses it, so it is not imposed. Section 3 of the paper, the tail factor, is a convention: the tail factor and the variance parameters attached to it are supplied by the actuary, so they are definitions with no theorem attached.

9.1 Definitions

Definition 9.1 (Weighted column sums). $S_k^{w,\alpha} = \sum_{i=0}^{n-k-2} w_{ik} C_{ik}^\alpha$ and $T_k^{w,\alpha} = \sum_{i=0}^{n-k-2} w_{ik} C_{ik}^\alpha F_{ik}$; `unitWeights` is $w_{ik} = 1$.

Definition 9.2 (Weighted development factor and variance estimator). $\hat{f}_k = T_k^{w,\alpha}/S_k^{w,\alpha}$ and $\hat{\sigma}_k^2 = \frac{1}{n-k-2} \sum_i w_{ik} C_{ik}^\alpha (F_{ik} - \hat{f}_k)^2$. As in the 1993 case, division by zero is zero, and the value at $k = n - 2$ is an extrapolation convention rather than an estimator.

Definition 9.3 (Weighted projection and ultimate). $\hat{C}_{ik} = C_{i,n-1-i} \prod_{j=n-1-i}^{k-1} \hat{f}_j$ and $\hat{C}_{i,n-1}$, both built from the weighted factors.

Definition 9.4 (Mack 1999, equation (*)).

$$\text{mse}W_i = \hat{C}_{i,n-1}^2 \sum_{k=n-1-i}^{n-2} \frac{\hat{\sigma}_k^2}{\hat{f}_k^2} \left(\frac{1}{w_{ik} \hat{C}_{ik}^\alpha} + \frac{1}{S_k^{w,\alpha}} \right).$$

Definition 9.5 (Mack 1999, the recursion below (*)). m steps of $V_{m+1} = \hat{C}_{ik}^2 (\hat{\sigma}_k^2 / (w_{ik} \hat{C}_{ik}^\alpha) + \hat{\sigma}_k^2 / S_k^{w,\alpha}) + V_m \hat{f}_k^2$ with $k = n - 1 - i + m$, started at $V_0 = 0$ on the latest observed diagonal.

Definition 9.6 (Tail factor, Section 3). The ultimate extended by a tail factor, $\hat{C}_{i,\infty} = \hat{C}_{i,n-1} f_{\text{tail}}$, and one further step of the recursion with actuary-supplied $f_{\text{tail}}, \sigma_{\text{tail}}^2, S_{\text{tail}}$. The paper gives no estimator for these, only the arithmetic of carrying them through; they are conventions and no theorem here uses them.

9.2 Deterministic identities

Lemma 9.7 (Weighted-average form). $\hat{f}_k = \sum_i (w_{ik} C_{ik}^\alpha / S_k^{w,\alpha}) F_{ik}$: the generalized factor is the $w_{ik} C_{ik}^\alpha$ -weighted mean of the individual factors. This is Lemma 2.9 with the weights carried through; because the numerator is written with the individual factors, no nonvanishing hypothesis is needed.

Proof. Divide the numerator termwise by $S_k^{w,\alpha}$. □

Lemma 9.8 (Weighted sum of squares). For any centre f , $\sum_i w_{ik} C_{ik}^\alpha (F_{ik} - f)^2 = \sum_i w_{ik} C_{ik}^\alpha F_{ik}^2 - 2f T_k^{w,\alpha} + f^2 S_k^{w,\alpha}$; at $f = \hat{f}_k$, and provided $S_k^{w,\alpha} \neq 0$, the cross term collapses to $\sum_i w_{ik} C_{ik}^\alpha F_{ik}^2 - S_k^{w,\alpha} \hat{f}_k^2$. These generalize Lemmas 2.11 and 2.12, the identity behind the $n - k - 2$ degrees of freedom.

Proof. Expand termwise; substitute $\hat{f}_k = T_k^{w,\alpha} / S_k^{w,\alpha}$. □

Lemma 9.9 (Weighted projection identities). $\hat{C}_{i,k+1} = \hat{C}_{ik} \hat{f}_k$ for $k \geq n - 1 - i$, $\hat{C}_{i,n-1-i} = C_{i,n-1-i}$, and (*) is the ultimate squared times the sum of its summands.

Proof. Split the product at the top; the empty product is one. □

9.3 The recursion equals (*)

Lemma 9.10 (One recursion step). For reals A, B, P, S, f, X with $f \neq 0$,

$$A^2 \left(\frac{B}{P} + \frac{B}{S} \right) + A^2 X f^2 = (A f)^2 \left(X + \frac{B}{f^2} \left(\frac{1}{P} + \frac{1}{S} \right) \right).$$

With $A = \hat{C}_{ik}$, $f = \hat{f}_k$, $B = \hat{\sigma}_k^2$, $P = w_{ik} \hat{C}_{ik}^\alpha$ and $S = S_k^{w,\alpha}$: one step of the recursion adds exactly the k -th summand of (*), scaled by the next projection. Only $\hat{f}_k \neq 0$ is used.

Proof. Clear the f^2 and expand. $\square$

Lemma 9.11 (Truncated form). *If $\hat{f}_k \neq 0$ for $k = n - 1 - i, \dots, n - 2 - i + m$, then m steps of the recursion equal $\hat{C}_{i,n-1-i+m}^2 \sum_{k=n-1-i}^{n-2-i+m}$ of the summands of $(*)$.*

Proof. Induction on m , using Lemma 9.10 at each step. $\square$

Theorem 9.12 (Mack 1999: the recursion is $(*)$, for general w and α). *Let $i \leq n - 1$ and suppose $\hat{f}_k \neq 0$ for $k = n - 1 - i, \dots, n - 2$. Then i steps of the recursion displayed below Mack's $(*)$ return $(*)$ itself, for every weight function and every natural-number exponent.*

Proof. Lemma 9.11 at $m = i$, where $n - 1 - i + i = n - 1$. $\square$

9.4 The 1993 case

Lemma 9.13 (Unit weights, $\alpha = 1$). *At $\alpha = 1$ with $w \equiv 1$, and provided the contributing entries C_{ik} do not vanish, $S_k^{w,\alpha} = S_k$, $T_k^{w,\alpha} = T_k$, $\hat{f}_k$ is the 1993 factor, $\hat{\sigma}_k^2$ is Mack's 1993 variance estimator, and the projection, the ultimate and the summands of $(*)$ are the 1993 ones.*

Proof. Termwise; $C_{ik}F_{ik} = C_{i,k+1}$ needs $C_{ik} \neq 0$. $\square$

Theorem 9.14 (Specialization to Mack 1993). *At $\alpha = 1$ with unit weights, $(*)$ is $\widehat{\text{mse}}(\hat{R}_i)$ and the generalized recursion is the recursion of Theorem 3.5; consequently i steps of it return Mack's 1993 estimator, which is the statement of Theorem 3.5.*

Proof. Rewrite with Lemma 9.13 and apply Theorem 9.12. $\square$

9.5 Conditional unbiasedness of the weighted estimator

Definition 9.15 (The weighted estimator as a random variable). $S_k^{w,\alpha}$ and $\hat{f}_k$ evaluated at the random triangle, with weights $w_{ik} : \Omega \rightarrow \mathbb{R}$, together with the multiplier $g_{ik} = w_{ik}C_{ik}^\alpha/C_{ik}$ of $C_{i,k+1}$ inside $\hat{f}_k$. For $\alpha = 1$ and unit weights $g_{ik} = 1$; writing the multiplier this way keeps $\alpha = 0$, where it is $1/C_{ik}$, in range.

Lemma 9.16 (Measurability and the linear form). *If every w_{ik} is $\mathcal{D}_k$ -measurable then so are $S_k^{w,\alpha}$ and g_{ik} (for $\alpha = 2$ this uses only measurability of C_{ik} , since C_{ik}^2 is then measurable), and $\hat{f}_k = (S_k^{w,\alpha})^{-1} \sum_i g_{ik}C_{i,k+1}$: the weighted estimator is a $\mathcal{D}_k$ -measurable linear combination of the next column.*

Proof. Sums, products, powers and quotients of measurable functions; the linear form is $w_{ik}C_{ik}^\alpha F_{ik} = (w_{ik}C_{ik}^\alpha/C_{ik})C_{i,k+1}$. $\square$

Theorem 9.17 (The weighted estimator is conditionally unbiased). *Under (M1) in the $\mathcal{D}_k$ form, with $\mathcal{D}_k$ -measurable weights, on the event that the weighted column sum is nonzero and where the contributing entries C_{ik} do not vanish, $E[\hat{f}_k \mid \mathcal{D}_k] = f_k$ for every natural-number exponent α . This is Theorem 5.5 with the weights and the exponent carried through.*

Proof. By Lemma 9.16 the factor $(S_k^{w,\alpha})^{-1}$ is $\mathcal{D}_k$ -measurable and comes out of the conditional expectation, and so does each g_{ik} . (M1) turns $E[C_{i,k+1} \mid \mathcal{D}_k]$ into $f_k C_{ik}$, and $g_{ik}C_{ik} = w_{ik}C_{ik}^\alpha$ when $C_{ik} \neq 0$, so the sum is $f_k S_k^{w,\alpha}$. The hypothesis $C_{ik} \neq 0$ is what the individual factors need in any case and cannot be dropped at $\alpha = 0$. $\square$

9.6 Weighted variance and estimator unbiasedness

Definition 9.18 (Weighted variance assumptions). Write $a_{ik} = w_{ik}C_{ik}^\alpha$ and $d_{ik} = F_{ik} - f_k$. The weighted conditional second-moment assumption is

$$E[d_{ik}^2 \mid \mathcal{D}_k] = \frac{\sigma_k^2}{a_{ik}}.$$

The cross-year assumption says $E[d_{ik}d_{jk} \mid \mathcal{D}_k] = 0$ for $i \neq j$. These are the $\mathcal{D}_k$ -conditioned forms used in the calculation.

Lemma 9.19 (Weighted residual algebra). *For predictable weights and $S_k^{w,\alpha} \neq 0$,*

$$\hat{f}_k^w - f_k = (S_k^{w,\alpha})^{-1} \sum_i a_{ik} d_{ik}.$$

Squaring gives the corresponding double sum. The weighted residual sum of squares also satisfies

$$\sum_i a_{ik} (F_{ik} - \hat{f}_k^w)^2 = \sum_i a_{ik} d_{ik}^2 - S_k^{w,\alpha} (\hat{f}_k^w - f_k)^2.$$

Proof. Expand the weighted sums and use $T_k^{w,\alpha} = \hat{f}_k^w S_k^{w,\alpha}$. □

Theorem 9.20 (Conditional variance of the weighted factor). *Under the weighted second-moment and cross-year assumptions, with predictable weights, nonzero contributing claims and weighted volumes, nonzero $S_k^{w,\alpha}$, and the stated integrability hypotheses,*

$$E[(\hat{f}_k^w - f_k)^2 \mid \mathcal{D}_k] = \frac{\sigma_k^2}{S_k^{w,\alpha}}.$$

When $E[\hat{f}_k^w \mid \mathcal{D}_k] = f_k$, the same expression is the conditional variance $\text{Var}(\hat{f}_k^w \mid \mathcal{D}_k)$.

Proof. Pull out the measurable weighted volumes, discard the cross terms, apply the weighted second-moment assumption on the diagonal, and cancel one weighted column sum. □

Definition 9.21 (Weighted variance estimator as a random variable). Evaluate $\hat{\sigma}_k^2$ and its weighted residual sum of squares at the random triangle and predictable weights.

Theorem 9.22 (Unbiasedness of the weighted variance estimator). *If $k + 3 \leq n$ and the hypotheses of Theorem 9.20 hold with the additional integrability needed by the weighted residual sum of squares, then*

$$E[\text{WSS}_k \mid \mathcal{D}_k] = (n - k - 2)\sigma_k^2, \quad E[\hat{\sigma}_k^2 \mid \mathcal{D}_k] = \sigma_k^2.$$

Every contributing a_{ik} is required to be nonzero. If zero-weight cells are omitted, an active-contributor set and its cardinality minus one must replace the fixed denominator.

Proof. The first term of the residual decomposition contributes $(n - k - 1)\sigma_k^2$ and Theorem 9.20 makes the correction term σ_k^2 . Divide their difference by $n - k - 2$. □

Chapter 10

The over-dispersed Poisson model

Renshaw and Verrall, *A stochastic model underlying the chain-ladder technique* [12], put the run-off triangle in incremental form, $X_{i,k} = C_{i,k} - C_{i,k-1}$ with $X_{i,0} = C_{i,0}$, and model the incrementals as an over-dispersed Poisson (ODP) generalized linear model: independent, with mean $m_{i,k}$ and variance $\varphi m_{i,k}$, and a log link carrying a row effect and a column effect, so the mean is multiplicative,

$$m_{i,k} = a_i b_k, \quad \sum_{k=0}^{n-1} b_k = 1.$$

Mack, *A simple parametric model for rating automobile insurance or estimating IBNR claims reserves* [13], starts from the same multiplicative mean and estimates it from the *marginal sums*. Both papers arrive at the same system: for each accident year i and each development year k ,

$$\sum_{k: i+k \leq n-1} m_{i,k} = \sum_{k: i+k \leq n-1} X_{i,k}, \quad \sum_{i: i+k \leq n-1} m_{i,k} = \sum_{i: i+k \leq n-1} X_{i,k}.$$

These are the score equations of the quasi-Poisson log-likelihood $\sum (X \log m - m)/\varphi$ in the row and column parameters; the dispersion φ multiplies the whole score and cancels, which is why the ODP fit is the Poisson fit and why the estimates do not depend on φ .

The chapter proves the two halves of the classical equivalence. Chain ladder *solves* the system (Theorems 10.10 and 10.12), and a positive multiplicative solution of the system *is* chain ladder (Theorem 10.16).

Source note. Neither paper was available in full text for this formalization. What is formalized is the marginal-sum system as quoted above and as it is stated in the secondary literature: England and Verrall [5, Section 2.3] and Wüthrich and Merz [14, Section 2.3]. The displayed equations of Renshaw and Verrall (Section 3) and of Mack (Section 2) are therefore identified by section and by the display itself, not by an equation number. Nothing here is probabilistic: the ODP assumption enters only through the shape of the score equations, which are algebraic identities in the fitted values.

10.1 Definitions

Definition 10.1 (Incremental claims). $X_{i,k} = C_{i,k} - C_{i,k-1}$ for $k \geq 1$ and $X_{i,0} = C_{i,0}$.

Definition 10.2 (Observed row and column). The cell (i, k) is observed when $i + k \leq n - 1$. The observed development years of accident year i are $\text{obsRow}(n, i) = \{0, \dots, n - 1 - i\}$ and the

observed accident years of development year k are $\text{obsCol}(n, k) = \{0, \dots, n-1-k\}$. The second one at $k+1$ is the contributor set of $\hat{f}_k$ (obsCol_succ).

Definition 10.3 (Chain-ladder fit on the observed triangle). The fitted cumulative on an observed cell is the latest observed entry of the row divided back by the development factors,

$$\hat{C}_{i,k} = \frac{C_{i,n-1-i}}{\prod_{j=k}^{n-2-i} \hat{f}_j},$$

and the fitted incremental is $\hat{X}_{i,k} = \hat{C}_{i,k} - \hat{C}_{i,k-1}$ with $\hat{X}_{i,0} = \hat{C}_{i,0}$. These are the $\hat{C}$ projections of Definition 2.6 read backwards from the latest diagonal: equivalently $\hat{C}_{i,k} = \hat{C}_{i,n-1} / \prod_{j=k}^{n-2-i} \hat{f}_j$, the chain-ladder ultimate times the cumulative development pattern up to k ($\text{CLcum_eq_ultimate_div}$).

Definition 10.4 (Multiplicative model and its pattern). $m_{i,k} = a_i b_k$; the identifiability constraint $\sum_{k < n} b_k = 1$; and the cumulative pattern $B_k = \sum_{j \leq k} b_j$.

Definition 10.5 (Score (marginal-sum) equations). The row marginal sums, the column marginal sums, and their conjunction: the fitted values reproduce the observed row and column totals of incrementals on the observed triangle.

Definition 10.6 (Quasi-Poisson log-likelihood). $\sum_k (X_{i,k} \log(t b_k) - t b_k)$ as a function of the row parameter t , and $\sum_i (X_{i,k} \log(a_i t) - a_i t)$ as a function of the column parameter t (dispersion $\varphi = 1$, which only scales the expression).

10.2 The score equations are score equations

Lemma 10.7 (Row and column scores). *The derivative of the quasi-Poisson log-likelihood in a row parameter is $(\sum_k X_{i,k} - \sum_k t b_k)/t$, and in a column parameter $(\sum_i X_{i,k} - \sum_i a_i t)/t$.*

Proof. Termwise: $\frac{d}{dt}(X \log(tb) - tb) = X/t - b$, then sum over the observed cells. $\square$

Lemma 10.8 (Stationarity is the marginal-sum equation). *For nonzero parameters, the row score vanishes exactly at the row marginal-sum equation and the column score exactly at the column one. This is what licenses the name "score equations" for Definition 10.5, and it is where the dispersion φ drops out.*

Proof. Lemma 10.7 and $a/b = 0 \iff a = 0$ for $b \neq 0$. $\square$

10.3 Chain ladder solves the score equations

Lemma 10.9 (Telescoping). $\sum_{k \leq m} X_{i,k} = C_{i,m}$ and $\sum_{k \leq m} \hat{X}_{i,k} = \hat{C}_{i,m}$.

Proof. Induction on m . $\square$

Theorem 10.10 (Row totals). *For every accident year i , $\sum_{k \leq n-1-i} \hat{X}_{i,k} = C_{i,n-1-i} = \sum_{k \leq n-1-i} X_{i,k}$: the chain-ladder fit reproduces every observed row total. This is the first marginal-sum family of Mack [13, Section 2], and the row score equation of Renshaw and Verrall [12, Section 3]. No hypothesis is needed.*

Proof. Both sides telescope to the latest observed entry, by Lemma 10.9 and $\hat{C}_{i,n-1-i} = C_{i,n-1-i}$ (the empty product). $\square$

Lemma 10.11 (Fitted column sums of cumulatives). *If every development factor $\hat{f}_j$, $j < n - 1$, is nonzero, then for every $k \leq n - 1$*

$$\sum_{i \leq n-1-k} \hat{C}_{i,k} = \sum_{i \leq n-1-k} C_{i,k},$$

and, one year earlier over the same accident years, $\sum_{i \leq n-2-k} \hat{C}_{i,k} = S_k$.

Proof. Downward induction on k . At $k = n - 1$ the only observed cell is the corner, where the fit is the observation. For the step, the cells strictly above the diagonal satisfy $\hat{C}_{i,k} = \hat{C}_{i,k+1}/\hat{f}_k$, so their sum is $T_k/\hat{f}_k = S_k$ by the induction hypothesis and the definition $\hat{f}_k = T_k/S_k$; adding the diagonal cell, where the fit is the observation, gives the claim. Nonvanishing of $\hat{f}_k$ forces $T_k \neq 0$ and $S_k \neq 0$, which is the only place a hypothesis is used. $\square$

Theorem 10.12 (Column totals). *If every development factor $\hat{f}_j$, $j < n - 1$, is nonzero, then for every development year k*

$$\sum_{i \leq n-1-k} \hat{X}_{i,k} = \sum_{i \leq n-1-k} X_{i,k} :$$

the chain-ladder fit reproduces every observed column total. This is the second marginal-sum family of Mack [13, Section 2], and the column score equation of Renshaw and Verrall [12, Section 3].

Proof. For $k = 0$ the incrementals are the cumulatives and this is Lemma 10.11. For $k \geq 1$ subtract the two displays of Lemma 10.11: the first over the observed part of column k , the second at $k - 1$ over the same accident years. $\square$

Theorem 10.13 (Chain ladder solves the ODP score equations). *Under the same hypothesis, the chain-ladder fitted incrementals satisfy the whole system of Definition 10.5.*

Proof. Theorems 10.10 and 10.12. $\square$

10.4 Uniqueness: the marginal-sum solution is chain ladder

Mack [13, Section 2] states the converse: the marginal-sum estimates of the multiplicative model are the chain-ladder estimates. Positivity of the parameters replaces the informal assumption that the triangle is a real one.

Lemma 10.14 (The pattern ratios are the development factors). *Let $a_i, b_k > 0$ with $\sum_{k < n} b_k = 1$ satisfy the score equations. Then for every $k \leq n - 1$*

$$\left(\sum_{i \leq n-1-k} a_i \right) B_k \text{ splits as } \sum_{i \leq n-1-k} C_{i,k}, \quad \text{and} \quad B_k \hat{f}_k = B_{k+1} \quad (k < n - 1),$$

so every $\hat{f}_k$ is positive.

Proof. Downward induction on k , mirroring Lemma 10.11. The row equation for accident year i reads $a_i B_{n-1-i} = C_{i,n-1-i}$; at $i = 0$, with $B_{n-1} = 1$, it starts the induction. In the step, write $A = \sum_{i \leq n-2-k} a_i$. The induction hypothesis gives $AB_{k+1} = T_k$ and the column equation at $k + 1$ gives $Ab_{k+1} = T_k - S_k$, so $AB_k = S_k$; adding the diagonal cell, whose row equation is $a_{n-1-k} B_k = C_{n-1-k,k}$, closes the induction. Dividing the two displays, $\hat{f}_k = T_k/S_k = B_{k+1}/B_k$. $\square$

Lemma 10.15 (Fitted cumulatives agree). *Under the hypotheses of Lemma 10.14, $a_i B_k = \hat{C}_{i,k}$ on every observed cell.*

Proof. Telescoping $B_k \prod_{j=k}^{d-1} \hat{f}_j = B_d$ with $d = n - 1 - i$, and the row equation $a_i B_d = C_{i,d}$; the product is nonzero because each factor is. $\square$

Theorem 10.16 (Mack's uniqueness result). *A multiplicative model $m_{i,k} = a_i b_k$ with $a_i, b_k > 0$ and $\sum_{k < n} b_k = 1$ that satisfies the marginal-sum equations coincides with the chain-ladder fit cell by cell on the observed triangle: $a_i b_k = \hat{X}_{i,k}$. With Theorem 10.13 this is the classical equivalence: the over-dispersed Poisson generalized linear model and the chain-ladder method give the same fitted values, hence the same development pattern and the same reserves.*

Proof. Differencing Lemma 10.15 at k and $k - 1$, since $B_k - B_{k-1} = b_k$. $\square$

Chapter 11

A finite independent-row witness

The passage from Mack's row-conditioned assumptions to the observed-data filtration is realized by a concrete model. This section supplies one finite probability space satisfying every hypothesis used by Theorems 5.30 and 5.31. The model is a repository witness, while the assumptions it realizes are Mack's equations (1) and (2) and the filtration B_k used in the proof of Theorem 2 [1].

Definition 11.1 (Eight-outcome triangle with row-generated filtration). There are three independent sign coordinates on eight equally likely outcomes. Claims in accident year i are constant at development year zero and depend only on sign ξ_i thereafter. The filtration at development year k is defined to be the join of the claims in every row through k . The positive-sign event in row i is called `shockEvent i`.

Theorem 11.2 (Independent rows and the generated filtration). *The three row shocks are mutually independent. Every claim in row i is measurable with respect to that row's shock, so the full row sigma-algebras are independent. By construction, the observed-data filtration is exactly the join of the row histories.*

Proof. The eight intersections of the three positive-sign events have the product probabilities 1, 1/2, 1/4 and 1/8 according to their cardinality. Independence passes to the smaller sigma-algebras generated by the claims. The filtration identity is definitional. $\square$

Theorem 11.3 (The independence passage is non-vacuous). *The model satisfies Mack's row-conditioned mean equation (1). The general independence theorems therefore give the corresponding mean equation conditioned on $\mathcal{D}_k$ and the cross-row residual condition (M2'). Two outcomes give different claims in the same accident year, so this is a genuinely random instance of both conclusions.*

Proof. At the first development step, conditional expectation over the trivial row history averages the two signs. Every later claim is the stated development factor times the previous claim and is already measurable in the row history. Apply Theorems 5.30 and 5.31; direct evaluation at outcomes zero and one proves nontriviality. $\square$

Chapter 12

Merz-Wüthrich: the MSEP of the one-year claims development result

Merz and Wüthrich, *Modelling the claims development result for solvency purposes* [7], quantify the uncertainty of the one-year move of the predicted ultimate. The previous chapter defined the true claims development result and proved that it has conditional mean zero; this chapter adds its second moment, which is the conditional mean squared error of predicting it by the budget value 0, and records the estimator the paper displays for the observable claims development result.

Index conventions. The paper writes I for the last accident year, $J = I$ for the last development year, and $\mathcal{D}_I$ for the triangle after I diagonals. Here $I = J = n - 1$; $I - i$ is $n - 1 - i$, $C_{I-j,j}$ is $C_{n-1-j,j}$, S_j^I is S_j computed with n accident years and S_j^{I+1} is the same column sum computed with $n + 1$, that is, with the next diagonal added. Both are functions of the triangle observed today, because $S_j^{I+1} = S_j^I + C_{I-j,j}$ and $C_{I-j,j}$ sits on the latest observed diagonal.

The step formalized below is the development-year step $\mathcal{D}_k \rightarrow \mathcal{D}_{k+1}$ of this library's filtration, not the paper's calendar-year step; the caveat is the one recorded for Theorem 7.4.

12.1 The exact one-year process variance

Definition 12.1 (One-year process variance). For claims c at development year k and ultimate development year N ,

$$\text{cdrProcVar}(c, f, \sigma^2, N, k) = \left(\prod_{j=k+1}^{N-1} f_j \right)^2 \sigma_k^2 c.$$

Theorem 12.2 (Merz-Wüthrich 2008, display (2.18)). *Under (M1) and (M3), for $i < n$ and $k < n - 1$ and integrable claims,*

$$E[\text{CDR}_i(k)^2 \mid \mathcal{D}_k] = \left(\prod_{j=k+1}^{n-2} f_j \right)^2 \sigma_k^2 C_{i,k}$$

almost surely. Since the true CDR has conditional mean zero, this is at once its conditional variance and the conditional mean squared error of predicting it by 0, which is the budget value the paper enters in the income statement for prior accident years.

Proof. By Theorem 7.4 the true CDR is $(\prod_{j>k} f_j)$ times the one-step residual $\varepsilon_{i,k} = C_{i,k+1} - f_k C_{i,k}$, up to sign. The constant comes out of the conditional expectation squared, and $E[\varepsilon_{i,k}^2 | \mathcal{D}_k] = \sigma_k^2 C_{i,k}$ is assumption (M3). The conditional-variance form subtracts Theorem 7.2 squared, which is 0. $\square$

12.2 The estimator of Merz and Wüthrich

Definitions only. The derivation in the paper's appendix replaces the estimated development factors by resampled ones and drops terms of higher order in the residuals, so none of the objects below is upgraded to an identity.

Lemma 12.3 (The column sum one year later). *For $j < n$, $S_j^{I+1} = S_j^I + C_{I-j,j}$.*

Proof. The index set gains exactly the row $I - j$. $\square$

Definition 12.4 (Displays (3.4)-(3.7)). With $d = I - i$,

$$\begin{aligned}\hat{\Psi}_i^I &= \frac{\hat{\sigma}_d^2 / (\hat{f}_d^I)^2}{C_{i,d}}, & \hat{\Phi}_{i,J}^I &= \sum_{j=d+1}^{J-1} \left(\frac{C_{I-j,j}}{S_j^{I+1}} \right)^2 \frac{\hat{\sigma}_j^2 / (\hat{f}_j^I)^2}{C_{I-j,j}}, \\ \hat{\Gamma}_{i,J}^I &= \hat{\Phi}_{i,J}^I + \hat{\Psi}_i^I, & \hat{\Delta}_{i,J}^I &= \frac{\hat{\sigma}_d^2 / (\hat{f}_d^I)^2}{S_d^I} + \sum_{j=d+1}^{J-1} \left(\frac{C_{I-j,j}}{S_j^{I+1}} \right)^2 \frac{\hat{\sigma}_j^2 / (\hat{f}_j^I)^2}{S_j^I}.\end{aligned}$$

The later development steps enter the one-year figure only through the re-estimation of their development factors on the next diagonal, which is why each is scaled by $(C_{I-j,j}/S_j^{I+1})^2$. No interval bound is claimed for general real-valued triangles.

Definition 12.5 (Displays (3.8)-(3.10)). $\widehat{\text{Var}}(\text{CDR}_i(I+1) | \mathcal{D}_I) = (\hat{C}_{i,J}^I)^2 \hat{\Psi}_i^I$; the prospective estimator $\widehat{\text{mse}}(0) = (\hat{C}_{i,J}^I)^2 (\hat{\Gamma}_{i,J}^I + \hat{\Delta}_{i,J}^I)$ split into its process part $(\hat{C}_{i,J}^I)^2 \hat{\Gamma}_{i,J}^I$ and its parameter part $(\hat{C}_{i,J}^I)^2 \hat{\Delta}_{i,J}^I$; and the retrospective estimator $\widehat{\text{mse}}(\widehat{\text{CDR}}_i(I+1)) = (\hat{C}_{i,J}^I)^2 (\hat{\Phi}_{i,J}^I + \hat{\Delta}_{i,J}^I)$.

Lemma 12.6 (Display (3.11)). $\widehat{\text{mse}}(0) = \widehat{\text{mse}}(\widehat{\text{CDR}}_i(I+1)) + \widehat{\text{Var}}(\text{CDR}_i(I+1) | \mathcal{D}_I)$, and the second is at most the first when $\hat{\Psi}_i^I \geq 0$.

Proof. $\hat{\Gamma} = \hat{\Phi} + \hat{\Psi}$. $\square$

Theorem 12.7 (Display (3.17)). *If the column sums S_j^I and S_j^{I+1} along the row do not vanish, then*

$$\widehat{\text{mse}}(0) = (\hat{C}_{i,J}^I)^2 \left(\frac{\hat{\sigma}_d^2 / (\hat{f}_d^I)^2}{C_{i,d}} + \frac{\hat{\sigma}_d^2 / (\hat{f}_d^I)^2}{S_d^I} + \sum_{j=d+1}^{J-1} \frac{C_{I-j,j}}{S_j^{I+1}} \frac{\hat{\sigma}_j^2 / (\hat{f}_j^I)^2}{S_j^I} \right).$$

Read next to Mack's Theorem 3, the one-year figure keeps only the first term of the process variance, keeps the estimation error of the next diagonal in full, and scales every remaining run-off cell by $C_{I-j,j}/S_j^{I+1}$. No interval bound is claimed without additional sign hypotheses.

Proof. The two sums combine term by term: with $s = S_j^I$, $c = C_{I-j,j}$ and $s + c = S_j^{I+1}$ by Lemma 12.3, $(c/(s+c))^2 a/c + (c/(s+c))^2 a/s = (c/(s+c)) a/s$. $\square$

12.3 The bridge to the exact statement

Theorem 12.8 (The first development step is the plug-in). *If the row is not fully developed and $\hat{f}_d^I \neq 0$, $C_{i,d} \neq 0$, then $(\hat{C}_{i,J}^I)^2 \hat{\Psi}_i^I$ is exactly $\text{cdrProcVar}(C_{i,d}, \hat{f}, \hat{\sigma}^2, n-1, d)$: display (3.8) is the plug-in of Theorem 12.2 for the step the next calendar year observes. The process part of display (3.9) is that plug-in plus the correction $(\hat{C}_{i,J}^I)^2 \hat{\Phi}_{i,J}^I$ for the later steps.*

Proof. Split the ultimate at its bottom factor, $\hat{C}_{i,J}^I = C_{i,d} \hat{f}_d^I \prod_{j>d} \hat{f}_j^I$; the squares of $C_{i,d}$ and $\hat{f}_d^I$ cancel against the denominators of $\hat{\Psi}_i^I$. $\square$

Theorem 12.9 (The one-year process variances telescope). *Summing the exact one-year process variance of Theorem 12.2 over every remaining development step of the row, with $C_{i,k}$ replaced by its (M1) prediction $C_{i,d} \prod_{l=d}^{k-1} \hat{f}_l$, gives Mack's multi-year process variance V_m exactly, with no remainder. The stochastic form performs the replacement by a conditional expectation given $\mathcal{D}_d$ instead of by hand.*

Proof. Term by term this is the closed form of V_m (Definition 5.20): its j -th summand is the one-year process variance of step $d+j$. For the stochastic form, each summand is a constant times $C_{i,d+j}$, and Lemma 7.3 evaluates its conditional expectation. $\square$

12.4 Aggregation over accident years

Accident year 0 is fully developed, so the sums run over $i = 1, \dots, n-1$ and the cross terms over the pairs $k > i > 0$, the older year i carrying the shared development factors. Definitions only.

Definition 12.10 (Displays (3.12)-(3.15)).

$$\hat{\Lambda}_{i,J}^I = \frac{C_{i,d}}{S_d^{I+1}} \frac{\hat{\sigma}_d^2 / (\hat{f}_d^I)^2}{S_d^I} + \sum_{j=d+1}^{J-1} \left(\frac{C_{i-j,j}}{S_j^{I+1}} \right)^2 \frac{\hat{\sigma}_j^2 / (\hat{f}_j^I)^2}{S_j^I}, \quad \hat{\Xi}_{i,J}^I = \hat{\Phi}_{i,J}^I + \frac{\hat{\sigma}_d^2 / (\hat{f}_d^I)^2}{S_d^{I+1}},$$

and the two aggregates $\sum_i \widehat{\text{mse}}_i + 2 \sum_{k>i>0} \hat{C}_{i,J}^I \hat{C}_{k,J}^I (\cdot_i + \hat{\Lambda}_{i,J}^I)$ with $\cdot_i = \hat{\Phi}_{i,J}^I$ for the retrospective view (3.12) and $\cdot_i = \hat{\Xi}_{i,J}^I$ for the prospective view (3.15).

Theorem 12.11 (Display (3.16)). *The aggregated prospective estimator differs from the aggregated retrospective one by $\sum_i \widehat{\text{Var}}(\text{CDR}_i(I+1) \mid \mathcal{D}_I) + 2 \sum_{k>i>0} \hat{C}_{i,J}^I \hat{C}_{k,J}^I (\hat{\Xi}_{i,J}^I - \hat{\Phi}_{i,J}^I)$: the same decoupling for aggregated accident years as for a single one.*

Proof. Termwise, Lemma 12.6 and ring algebra. $\square$

Chapter 13

The observed data: Mack's hypotheses discharged

13.1 From independence across accident years to the hypotheses

Theorem 5.23 conditions on a σ -algebra $\mathcal{D}$ of observed data and carries two named hypotheses: that $\mathcal{D}$ says nothing about accident year i beyond that year's own history up to its latest observed development year. Theorem 8.17 carries a third, Definition 8.12. Mack obtains all three from assumption (M2) in the sentence “because of the independence of the accident years” in the proof of Theorem 3 [1, pp. 217–219]. This section names the σ -algebra and proves that sentence.

Definition 13.1 (The observed data). After n calendar years the triangle holds one diagonal per accident year: row i is observed up to development year $n - 1 - i$. So

$$\text{obsSigma} = \bigvee_{i < n} \sigma(C_{i,0}, \dots, C_{i,n-1-i}),$$

and otherObs_i is the same join over the accident years $i' \neq i$.

Lemma 13.2 (The observed data splits at one accident year). *For $i < n$, $\text{obsSigma} = \sigma(C_{i,0}, \dots, C_{i,n-1-i}) \vee \text{otherObs}_i$, and otherObs_i is contained in the σ -algebra of all the other accident years.*

Proof. Both inclusions termwise: a term of the join either is row i 's or is one of the others'. This is Lemma 5.29 for the latest diagonal instead of a development year. $\square$

Lemma 13.3 (The chain-ladder estimators are functions of the observed data). *S_k , $\hat{f}_k$ and $\hat{C}_{i,d+m}$ are obsSigma -measurable, with no hypothesis: an accident year contributes to $\hat{f}_k$ only when $C_{i,k+1}$ lies on or above the latest diagonal. This is the hypothesis $\hat{C}_{i,d+m}$ is $\mathcal{D}$ -measurable of Theorem 5.23.*

Proof. $i \in \text{contributors}(n, k)$ gives $i + k + 2 \leq n$, hence $k + 1 \leq n - 1 - i$; then induction on m for the product. $\square$

Theorem 13.4 (No further information about one accident year). *Assume (M2) and that $\mathcal{D}_k$ is generated by the row histories. For g integrable and measurable for accident year i , and $i < n$,*

$$E[g \mid \text{obsSigma}] = E[g \mid \sigma(C_{i,0}, \dots, C_{i,n-1-i})] = E[g \mid \mathcal{D}_{n-1-i}].$$

Taking $g = C_{i,k}$ and $g = C_{i,k}^2$ gives exactly the two “no extra information” hypotheses of Theorem 5.23.

Proof. Split the conditioning σ -algebra by Lemma 13.2 (for the first equality) and by Lemma 5.29 (for the second), and apply Theorem 5.28 with m_1 row i ’s observed history, m'_1 the whole of row i and m_2 the other rows, which (M2) makes independent of m'_1 . Both sides reduce to the same conditional expectation. $\square$

Theorem 13.5 (Mack’s Theorem 3 from Mack’s own assumptions). *Assume (M1row), (M3row), (M2), that $\mathcal{D}_k$ is generated by the rows’ histories, and square integrability of the claims and of the chain-ladder prediction. Then, conditionally on the observed data,*

$$E[(\hat{C}_{i,d+m} - C_{i,d+m})^2 \mid \text{obsSigma}] = V_m + (\hat{C}_{i,d+m} - C_{i,d} \prod_{l < m} f_{d+l})^2, \quad d = n - 1 - i.$$

Nothing is left over: the two hypotheses of Theorem 5.23 are Theorem 13.4, the measurability of the predictor is Lemma 13.3, and (M1), (M3) in the $\mathcal{D}_k$ form are Theorem 5.30.

Proof. Substitute those four into Theorem 5.23. $\square$

Theorem 13.6 (The cross-term condition from independence). *Assume (M2) and square integrability of the claims. For $i \neq j$ with $i < n$,*

$$E[(C_i - E[C_i \mid \text{obsSigma}])(C_j - E[C_j \mid \text{obsSigma}]) \mid \text{obsSigma}] = 0,$$

so Definition 8.12 holds for the true ultimates over all n accident years.

Proof. The argument of Theorem 5.31. Condition on the observed data together with the whole of accident year j : row j ’s centred ultimate is measurable there and comes out of the conditional expectation. What remains is row i ’s centred ultimate conditioned on a σ -algebra that adds nothing about row i (Theorem 5.28), and the centring was by exactly that conditional expectation (Theorem 13.4), so it vanishes. The tower property brings the identity back to the observed data. $\square$

Theorem 13.7 (Mack’s Corollary from Mack’s own assumptions). *Under the same hypotheses as Theorem 13.5, taken over all n accident years,*

$$E\left[\left(\sum_i \hat{C}_i - \sum_i C_i\right)^2 \mid \text{obsSigma}\right] = \sum_i \text{procVar}_i + \left(\sum_i (\hat{C}_i - M_i)\right)^2 \geq \sum_i \text{procVar}_i.$$

The chain from Mack’s row-conditioned assumptions plus independence across accident years to the exact Theorem 3 and its total-reserve form is closed.

Proof. Theorem 8.17 with $\mathcal{D} = \text{obsSigma}$, its cross-term hypothesis supplied by Theorem 13.6 and its remaining hypotheses by Theorem 13.4 and Lemma 13.3. $\square$

Chapter 14

Mack's published numerical example

Definition 14.1 (Taylor-Ashe cumulative triangle). The 55 observed cumulative claims in Mack's Table 1 are represented exactly as integers. Unobserved cells are zero and do not enter any estimator.

Theorem 14.2 (Published development factors). *The nine development factors computed from Table 1 equal exact rational numbers whose decimal roundings are the values printed by [1].*

Proof. Each equality follows by normalization of the finite sums and division. □

Theorem 14.3 (Published estimable variances). *The eight variance values supplied by Mack's variance estimator equal exact rationals with the printed decimal roundings. The ninth value is specified by an extrapolation convention and is not asserted here as an estimator value.*

Proof. Each equality follows by normalization of the weighted squared deviations. □

Definition 14.4 (Total reserve for Table 1). The total reserve is the sum of the ten chain-ladder reserves computed from the published triangle.

Theorem 14.5 (Published total reserve). *The exact total reserve is strictly between 18,680.5 and 18,681.5 thousand, so it rounds to the 18,681 thousand printed in Mack's Table 2.*

Proof. Normalization of the ten reserve terms proves both strict inequalities. □

Chapter 15

Mack's model diagnostics

The source for this chapter is Thomas Mack, “Measuring the Variability of Chain Ladder Reserve Estimates” [17], Appendices G and H, pages 157–166. The rank and label distributions used below are extra null models, not consequences of the chain-ladder moment assumptions.

15.1 Development-factor correlation

Definition 15.1 (Adjacent-column rank statistic). For each adjacent pair of development-factor columns, restrict the earlier column to the accident years present in both columns, rerank both restricted columns, and compute Spearman's coefficient. Mack's triangle-wide statistic is the inverse-null-variance weighted average of these coefficients.

Definition 15.2 (Exact permutation null). For two complete, tie-free rank permutations, ρ is their centered-rank correlation. Its exact null fixes one rank vector and draws the other uniformly from all permutations.

Theorem 15.3 (Equation (G4) and its exact null moments). *For $q \geq 2$, the centered-rank correlation is*

$$1 - \frac{6 \sum_i (r_i - s_i)^2}{q^3 - q}.$$

Under the uniform-permutation null its mean is zero and its variance is $1/(q - 1)$.

Proof. Reindex finite sums by permutations, center the ranks, and evaluate the diagonal and off-diagonal second moments by exchangeability. $\square$

Definition 15.4 (Nominal aggregate calibration). The inverse-variance aggregate has Mack's variance target only if its component coefficients are pairwise uncorrelated. The 0.67 Normal cutoff is recorded as a definition, not as an exact rejection-probability theorem.

15.2 Calendar-year labels

Definition 15.5 (Median split and diagonal score). Each development-factor column is split into small and large ranks around its median. A calendar diagonal scores the smaller of its two label counts, and the test statistic sums those scores.

Definition 15.6 (Exact fair-label null). For n retained labels, the exact null assigns binomial mass $\binom{n}{l}/2^n$ to l large labels and scores $\min(l, n - l)$. The finite mean, second moment, and variance are defined before their closed forms.

Theorem 15.7 (Calendar equations (H1) and (H2)). *The fair-label masses sum to one. For every n , the finite binomial sums equal Mack's closed mean and variance formulas (H1) and (H2); the second moment also has an explicit closed form.*

Proof. Pair symmetric binomial terms, telescope the signed half-sum, and evaluate the first two binomial factorial moments. $\square$

Theorem 15.8 (Five-label check). *For Mack's five-label example the exact mean, second moment, and variance are $50/32$, $90/32$, and $95/256$, and the closed formulas reproduce them.*

Proof. Normalize the finite binomial sums and evaluate them. $\square$

Definition 15.9 (Nominal calendar aggregation). Mack sums the component means and variances and applies a Normal cutoff. The variance sum and cutoff are nominal definitions because separate diagonal scores are described as almost uncorrelated rather than proved independent.

Chapter 16

Mack's lognormal approximation

The parameter conversion appears in Thomas Mack, “Measuring the Variability of Chain Ladder Reserve Estimates” [17], Chapter 4, equation (10), pages 118–119. The 1993 ASTIN paper does not contain this construction.

Definition 16.1 (Moment-matched parameters). For an arithmetic mean m and variance v , define

$$\sigma^2 = \log(1 + v/m^2), \quad \mu = \log m - \sigma^2/2.$$

These are the log-variance and log-mean parameters in Mack's equation (10).

Theorem 16.2 (Matched lognormal moments). *If $m > 0$ and $v \geq 0$, then*

$$\exp(\mu + \sigma^2/2) = m$$

and

$$\exp(2\mu + \sigma^2)(\exp(\sigma^2) - 1) = v.$$

Thus the standard lognormal moment expressions reproduce the supplied mean and variance. This identity does not state that the underlying reserve follows a lognormal law.

Proof. Substitute the parameter definitions, use $\exp(\log x) = x$ for positive x , and simplify the resulting rational identity. $\square$

Chapter 17

Prior uncertainty in Bornhuetter-Ferguson

Mack, *The Prediction Error of Bornhuetter/Ferguson* [9], calls the selected pattern and the a priori ultimate practically independent in Section 4, property (b), pages 94–95, then uses the product-variance formula on pages 95–96. This chapter freezes the observed chain-ladder pattern of Definition 4.1 and isolates the exact prior contribution. Independence below strengthens the paper’s description.

Definition 17.1 (Random BF reserve). For observed triangle X , accident year i , and random a priori ultimate U ,

$$A_i = 1 - \frac{1}{\text{cdf}_i(X)}, \quad R_i^{BF,U} = A_i U.$$

The coefficient A_i is measurable for the observed-data sigma-algebra.

Definition 17.2 (Independent prior). The sigma-algebra generated by U is independent of the observed-data sigma-algebra. The variance below is mathlib’s standard $\text{Var}(U)$, equal on a probability space to $\int (U - \int U d\mu)^2 d\mu$ under the stated measurability hypothesis.

Theorem 17.3 (Conditional mean of the random-prior BF reserve). *On a probability space, if U is strongly measurable and integrable, $A_i U$ is integrable, and the prior is independent of the observed triangle, then*

$$E[R_i^{BF,U} \mid \text{obsSigma}] = A_i E[U]$$

almost surely. This is the conditional form of the expectation calculation on Mack’s page 95.

Proof. Pull the observed coefficient out of the conditional expectation, then condition the independent prior on the observed-data sigma-algebra. $\square$

Theorem 17.4 (Conditional variance contributed by the prior). *If the prior and the resulting BF reserve are square-integrable, then*

$$E \left[(R_i^{BF,U} - A_i E[U])^2 \mid \text{obsSigma} \right] = A_i^2 \text{Var}(U)$$

almost surely. This is the prior-uncertainty component of the product-variance formula on Mack’s pages 95–96; it does not state the other components of the paper’s prediction-error approximation.

Proof. Rewrite the centred reserve as $A_i(U - E[U])$, pull A_i^2 out of the conditional expectation, and use independence for the squared centred prior. $\square$

Chapter 18

Bühlmann-Straub credibility

Bühlmann and Straub, *Glaubwürdigkeit für Schadensätze* [8], minimize expected squared prediction error over unbiased linear estimators in Section 6, pages 118–120, and derive the normal equations in display (1). Display (4), page 123, gives the exposure-weighted individual mean, the credibility factor, and the credibility-weighted blend. The deterministic loss below is an algebraic representation of the one-risk normal equation, not the expected prediction-error objective of the paper.

Definition 18.1 (Weighted experience, credibility, estimate, and loss). For a finite sample with exposure weights P_l and observations X_l , set

$$P_{\bullet} = \sum_l P_l, \quad T = \sum_l P_l X_l, \quad \bar{X} = T/P_{\bullet}, \quad Z = \frac{P_{\bullet} w}{v + P_{\bullet} w}.$$

For collective mean m , process variance v , and between-risk variance w , define

$$\hat{m} = \frac{wT + vm}{v + P_{\bullet} w}, \quad L(t) = w \sum_l P_l (X_l - t)^2 + v(t - m)^2.$$

The quantities m , v , and w are inputs.

Theorem 18.2 (Normal equation, square completion, and minimizer). *If $v + P_{\bullet} w \neq 0$, then*

$$(v + P_{\bullet} w)\hat{m} = wT + vm, \quad L(t) = L(\hat{m}) + (v + P_{\bullet} w)(t - \hat{m})^2.$$

Consequently, $L(\hat{m}) \leq L(t)$ for every t whenever $v + P_{\bullet} w > 0$.

Proof. Expand the weighted squared deviations, substitute the normal equation, and use non-negativity of a square. $\square$

Theorem 18.3 (Credibility form, display (4)). *If $P_{\bullet} \neq 0$ and $v + P_{\bullet} w \neq 0$, then*

$$\hat{m} = Z\bar{X} + (1 - Z)m.$$

If $P_{\bullet}$, v , and w are nonnegative and the denominator is positive, then $0 \leq Z \leq 1$.

Proof. Clear the two nonzero denominators and normalize the resulting polynomial identity. The interval follows by comparing the nonnegative numerator with its denominator. $\square$

Chapter 19

Munich chain ladder: the separate-chain-ladder gap

Quarg and Mack's Section 1.1.2 identifies the paid/incurred-ratio gap left when the two triangles are completed separately [3].

Definition 19.1 (Separate-chain-ladder quadrangle). The completed quadrangle keeps $C_{i,k}$ when $i + k \leq n - 1$ and uses $\hat{C}_{i,k}$ otherwise. Its column total includes all n rows.

Definition 19.2 (Paid/incurred ratios). For completed quadrangles P and I , $(P/I)_{i,k} = P_{i,k}/I_{i,k}$ and $(P/I)_k = (\sum_i P_{i,k})/(\sum_i I_{i,k})$.

Lemma 19.3 (Completed-column step). *If $k < n - 1$ and $S_k \neq 0$, the completed column total evolves by the original chain-ladder factor:*

$$\sum_i \hat{C}_{i,k+1} = \left(\sum_i \hat{C}_{i,k} \right) \hat{f}_k.$$

Proof. Split the observed prefix from the projected tail; use $T_k = \hat{f}_k S_k$ on the first and the projection step on the second. $\square$

Lemma 19.4 (Completed-column product). *For $d \leq t < n$, if the intervening factor denominators S_k are nonzero, then*

$$\sum_i \hat{C}_{i,t} = \left(\sum_i \hat{C}_{i,d} \right) \prod_{k=d}^{t-1} \hat{f}_k.$$

Proof. Induction on $t - d$ using Lemma 19.3. $\square$

Theorem 19.5 (Quarg–Mack gap identity). *For $i < n$ and $d = n - 1 - i < t < n$, assume the paid and incurred factors on $[d, t)$, $I_{i,d}$, and both completed column totals at d are nonzero. Then*

$$\frac{(P/I)_{i,t}}{(P/I)_t} = \frac{(P/I)_{i,d}}{(P/I)_d}.$$

Here $(P/I)_k$ uses every completed row. This is not a convergence claim and does not include the later Munich correlation adjustment.

Proof. Both the row values and the completed column totals acquire the same paid and incurred factor products; cancel them in the ratio of ratios. $\square$

Chapter 20

Munich residual correlations and adjusted factors

Quarg and Mack's Sections 2.2 and 3.1.2 add the stochastic link between paid and incurred development [3]. The model assumptions and the raw practical estimators are kept separate here.

Definition 20.1 (Joint row history and conditional residual). For paid and incurred random triangles on one probability space, $\mathcal{B}_i(k)$ is the join of their row histories through k . For a random variable Y and a smaller history $\mathcal{A}$ its residual is

$$\text{Res}(Y \mid \mathcal{A}) = \frac{Y - E[Y \mid \mathcal{A}]}{\sqrt{\text{Var}(Y \mid \mathcal{A})}}.$$

The PIU predicate states independence of the finite n accident years after grouping each year's paid and incurred histories through the final modeled period $n - 1$.

Definition 20.2 (PQ and IQ). PQ regresses the paid development residual, conditioned on $\mathcal{B}_i(k)$, on the reciprocal-ratio residual with one common slope λ^P . IQ uses the incurred development residual and the paid/incurred-ratio residual with slope λ^I . Both predicates range over adjacent periods satisfying $k + 1 < n$.

Theorem 20.3 (Residual slope is the standardized cross-moment). *Let $\mathcal{A} \leq \mathcal{B}$. If A is $\mathcal{B}$ -measurable, $E[Z \mid \mathcal{B}] = \lambda A$, and $E[A^2 \mid \mathcal{A}] = 1$, then*

$$E[AZ \mid \mathcal{A}] = \lambda.$$

Applying this to PQ and IQ proves the paid and incurred residual cross-moments equal λ^P and λ^I , under the stated integrability and standardization hypotheses.

Proof. Pull A out after conditioning on $\mathcal{B}$, substitute the regression, use the tower property, and use $E[A^2 \mid \mathcal{A}] = 1$. $\square$

Definition 20.4 (Raw Munich estimators and recursion). The raw pattern and dispersion estimators on pp. 290–291 define the four residual arrays. The pooled least-squares slopes through the origin are $\hat{\lambda}^P$ and $\hat{\lambda}^I$. At a projected pair (p, i) the paid factor is

$$\hat{f}_k^P + \hat{\lambda}^P \frac{\hat{\sigma}_k^P}{\hat{\rho}_k^P} \left(\frac{i}{p} - \hat{q}_k^{-1} \right),$$

with the symmetric paid/incurred formula for the incurred factor. The coupled recursion applies both factors to the preceding pair.

Theorem 20.5 (Zero residual slopes recover separate chain ladder). *When both slopes are zero, every adjusted factor is its usual chain-ladder factor and the coupled recursion is exactly $(\hat{P}_{i,d+m}, \hat{I}_{i,d+m})$ from the two separate projections.*

Proof. Induction on m , using the chain-ladder projection step for each component. □

Chapter 21

The stochastic ODP layer

England and Verrall describe independent incremental claims with

$$\mathbb{E}[X_{i,k}] = m_{i,k}, \quad \text{Var}(X_{i,k}) = \varphi m_{i,k}, \quad \log m_{i,k} = c + \alpha_i + \beta_k,$$

with $\varphi > 0$ and the corner constraints $\alpha_0 = \beta_0 = 0$ [4, Section 2, equations (2.5)–(2.7), p. 283]. The same moment and link equations appear in [6, Section 2, equations (2.1)–(2.3)]. This chapter separates exact consequences of those assumptions from the first-order GLM prediction-error formula and the residual bootstrap algorithm.

21.1 Model and conditional moments

Definition 21.1 (Finite stochastic ODP model). For a finite square of cells, the dispersion is positive and the mean surface admits a corner-constrained log-link representation. Every incremental claim is measurable and square-integrable, has that mean $m_{i,k}$ and variance $\varphi m_{i,k}$, and the family is mutually independent. In particular, every modeled mean is positive. The observed-data sigma-algebra is generated by any selected finite set of cells.

Lemma 21.2 (Generated sigma-algebras and independence). *The sigma-algebras generated by one cell and by a finite vector of cells are sub-sigma-algebras of the ambient one. A cell outside the selected vector is independent of that vector.*

Proof. Measurability gives the two inclusions. Mutual independence, restricted to the disjoint singleton and observed index sets, gives the final statement. $\square$

Theorem 21.3 (Exact conditional ODP moments). *If (i, k) is not among the observed cells, then*

$$\mathbb{E}[X_{i,k} \mid \mathcal{O}] = m_{i,k}, \quad \mathbb{E}[(X_{i,k} - m_{i,k})^2 \mid \mathcal{O}] = \varphi m_{i,k}$$

almost surely. These are exact conditional-moment identities, not GLM approximations.

Proof. Independence makes each conditional expectation equal to its integral; substitute the two fields of the ODP model. $\square$

21.2 Prediction error

Definition 21.4 (Aggregate and mean squared prediction error). For a finite set F of future cells, let $Y_F = \sum_{p \in F} X_p$. For a predictor $\hat{Y}$, define $\text{MSEP}(Y_F, \hat{Y}) = \mathbb{E}[(Y_F - \hat{Y})^2]$.

Theorem 21.5 (Aggregate mean and process variance). *For every finite F ,*

$$\mathbb{E}[Y_F] = \sum_{p \in F} m_p, \quad \text{Var}(Y_F) = \varphi \sum_{p \in F} m_p.$$

Proof. Integrate the finite sum and add the variances of the independent square-integrable cells. $\square$

Theorem 21.6 (Exact process-plus-estimation decomposition). *If a square-integrable predictor has the same mean as its square-integrable target and is independent of it, then*

$$\mathbb{E}[(Y - \hat{Y})^2] = \text{Var}(Y) + \text{Var}(\hat{Y}).$$

For an ODP aggregate this is $\varphi \sum_{p \in F} m_p + \text{Var}(\hat{Y})$.

Proof. The difference has mean zero. Expand its variance and use independence to set the covariance to zero, then substitute Theorem 21.5. $\square$

England and Verrall use the delta method for the estimation component: since $\hat{m} = \exp(\hat{\eta})$, $\text{Var}(\hat{m}) \doteq \hat{m}^2 \text{Var}(\hat{\eta})$. Their single-cell and aggregate formulas are equations (3.3)–(3.5) of [4] and equations (7.7)–(7.9) of [5]. The dotted equality is not promoted to a theorem.

Definition 21.7 (Delta-method GLM MSEP formulas). The single-cell formula is $\varphi \hat{m} + \hat{m}^2 \text{Var}(\hat{\eta})$. For a finite aggregate, the estimation term is the covariance quadratic form $\sum_{p,q} \hat{m}_p \hat{m}_q \text{Cov}(\hat{\eta}_p, \hat{\eta}_q)$, added to $\varphi \sum_p \hat{m}_p$. These are definitions of the source approximation.

Lemma 21.8 (Singleton bridge). *On a singleton future set, the aggregate definition is the single-cell formula, with the diagonal covariance used as $\text{Var}(\hat{\eta})$.*

Proof. Evaluate both finite sums on a singleton. $\square$

21.3 Log-link and corner constraints

Definition 21.9 (Log-link and corner effects). The log-link fit is $\exp(c + \alpha_i + \beta_k)$. For positive multiplicative parameters a_i, b_k , set $\alpha_i = \log a_i - \log a_0$ and $\beta_k = \log b_k - \log b_0$.

Theorem 21.10 (Log-link and multiplicative equivalence). *Every log-link fit is multiplicative. Conversely, every positive multiplicative fit has the corner-constrained representation above, with $c = \log a_0 + \log b_0$ and $\alpha_0 = \beta_0 = 0$.*

Proof. Use $\exp(x + y) = \exp(x) \exp(y)$ in one direction and $\exp(\log x) = x$ for positive x in the other. $\square$

21.4 Residual bootstrap definitions

England and Verrall’s unscaled Pearson residual, scale estimate, and inversion are equations (4.2)–(4.4) of [4]. England’s later two-stage procedure first resamples adjusted residuals and refits the mean, then simulates future process outcomes conditional on that mean [6, Sections 2–3].

Definition 21.11 (Pearson residual operations). For an observation x and fitted mean m , the unscaled Pearson residual is $r = (x - m)/\sqrt{m}$. The Pearson scale is $\sum r^2/(N - p)$, the degrees-of-freedom adjustment multiplies a residual by $\sqrt{N/(N - p)}$, and inversion gives the pseudo-observation $x^* = r^*\sqrt{m} + m$. Lean’s total division leaves these definitions meaningful at zero; the inversion theorems below assume $m > 0$.

Lemma 21.12 (Pearson inversion). *For $m > 0$, taking the residual of $r\sqrt{m} + m$ returns r , and rebuilding an observation from its Pearson residual returns that observation.*

Proof. Cancel the nonzero factor $\sqrt{m}$. □

Definition 21.13 (England-Verrall bootstrap and predictive reserve). One bootstrap refit resamples supplied residuals with replacement, inverts them against the fitted past means, and applies a supplied refitting map to obtain future means. The predictive reserve sums supplied process draws conditional on those bootstrap means. The process simulator is an input: the ODP moment equations do not determine a unique distribution. These declarations define the algorithm and make no coverage or consistency claim.

21.5 Finite non-vacuity witness

Theorem 21.14 (A nondegenerate ODP model exists). *On the product of four uniform two-point spaces, one coordinate is assigned to each cell of a 2 by 2 square. Every claim is 0 or 2, has mean and variance one, and the four claims are mutually independent. The unit mean surface satisfies the corner-constrained log link with zero effects and $\varphi = 1$. Taking $(0, 1)$ as future and the nonempty singleton $\{(0, 0)\}$ as observed instantiates both conditional-moment theorems. Two random coordinates also instantiate the exact MSEP decomposition with value 2.*

Proof. The coordinate projections are independent under the finite product measure. Transport the two-point mean and variance along each measure-preserving evaluation map. □

Chapter 22

Mack 1999: row assumptions and active weights

Mack's weighted CL2 is conditioned on the history of one accident year. The full-filtration formula and the cross-accident-year cancellation therefore come from the same independence step as their Mack 1993 counterparts. Zero weights are handled by fixing the active observations before fitting the weighted mean.

Definition 22.1 (Row-conditioned weighted CL2). The weights are measurable for the accident year's history and

$$\mathbb{E}[(F_{i,k} - f_k)^2 \mid C_{i,0}, \dots, C_{i,k}] = \frac{\sigma_k^2}{w_{i,k} C_{i,k}^\alpha}.$$

Theorem 22.2 (Transport of weighted CL2). *If the rows are independent and generate the observed-data filtration, then row-conditioned weighted CL2 implies Mack3W, its form conditioned on the full data through development year k .*

Proof. Split the filtration into row i and the other rows. Conditioning on the independent enlargement leaves the squared factor residual unchanged. $\square$

Theorem 22.3 (Cross-factor terms from independent rows). *Under row-conditioned CL1, nonzero cumulative claims and integrability, row independence gives*

$$\mathbb{E}[(F_{i,k} - f_k)(F_{j,k} - f_k) \mid D_k] = 0, \quad i \neq j.$$

Proof. Enlarge D_k by the whole of row j , take its residual outside the conditional expectation, use the zero conditional mean for row i , and apply the tower property. $\square$

Definition 22.4 (Active contributors and variance estimator). The active set is the fixed subset of contributors with nonzero deterministic weight. For an active set A ,

$$\hat{\sigma}_{k,A}^2 = \frac{1}{|A| - 1} \sum_{i \in A} w_{i,k} C_{i,k}^\alpha (F_{i,k} - \hat{f}_{k,A})^2.$$

When cumulative claims are nonzero, membership in A is equivalent to a nonzero weighted volume. Weighted CL2 on A imposes no condition on cells excluded from A .

Theorem 22.5 (Active weighted variance is unbiased). *For at least two active observations, weighted CL2, the cross-factor condition, nonzero active cumulative claims and weighted volumes, and the stated integrability hypotheses imply*

$$\mathbb{E}[\widehat{\sigma}_{k,A}^2 \mid D_k] = \sigma_k^2.$$

The set A is fixed; the theorem makes no claim for a selection rule that depends on the outcome.

Theorem 22.6 (Finite witnesses for $\alpha = 0$ and $\alpha = 2$). *The eight-outcome independent-row model satisfies the row-conditioned and full-filtration weighted assumptions, the derived cross-factor condition and active-estimator unbiasedness for both named exponents. Its first-step factor variance is positive and its claims are genuinely random.*

Chapter 23

Calendar-year claims development results

Merz and Wüthrich condition on the cells available after a calendar diagonal [7]. This chapter gives that filtration directly and keeps their Appendix A.1 linear approximation separate from the exact probability statements.

Definition 23.1 (Calendar-year information). At calendar time t ,

$$\mathcal{D}_t^{\text{cal}} = \sigma\{C_{i,j} : i < n, j < n, i + j \leq t\}.$$

The companion sigma-algebra `otherCalendar` contains the observed cells from every row except a specified accident year.

Theorem 23.2 (The calendar triangle splits at one row). *For $i < n$ and $k < n$, at time $i + k$ the calendar triangle is row i through development k joined with the observed parts of the other rows. Under accident-year independence, conditioning a row- i random variable on this calendar triangle is the same as conditioning it on row i through k .*

Proof. Split the generating joins at row i , then apply conditioning on an independent enlargement. $\square$

Definition 23.3 (True calendar-year CDR). The true claims development result over $t \rightarrow t + 1$ is

$$E[C_{i,n-1} \mid \mathcal{D}_t^{\text{cal}}] - E[C_{i,n-1} \mid \mathcal{D}_{t+1}^{\text{cal}}].$$

Theorem 23.4 (Exact calendar-year CDR). *The raw mean-zero equality needs only a finite measure; since conditional expectation is zero off the integrable case in the formalization, it has its intended probabilistic meaning when the ultimate is integrable. For $i < n$ and $k < n - 1$, under the row-conditioned Mack assumptions, independent accident years, row integrability, and squared-residual integrability, at time $i + k$ it equals*

$$\left(\prod_{j=k+1}^{n-2} f_j \right) (f_k C_{i,k} - C_{i,k+1}),$$

and its conditional second moment and conditional variance are $\left(\prod_{j=k+1}^{n-2} f_j \right)^2 \sigma_k^2 C_{i,k}$.

Proof. The tower property gives the mean-zero result. The split theorem reduces the other two statements to the row filtration and (M1row), (M3row). $\square$

Definition 23.5 (Appendix A.1 boundary). For finite increments a_j , the source replaces $\prod_j (1 + a_j) - 1$ by the first-order expression $\sum_j a_j$. Lean records the exact product, the linear approximation, and their difference as three separate definitions. No equality or statistical error bound is built into the approximation.

Theorem 23.6 (Exact boundary identities). *The product increment is the linear approximation plus its exact remainder. For nonnegative increments the linear term is a lower bound, as stated after (A.1), and for two increments the discarded term is exactly ab.*

Proof. Ring algebra and the finite Weierstrass product inequality. $\square$

Theorem 23.7 (Finite CDR and MSEP witness). *The eight-outcome independent-row model satisfies (M3row), instantiates the exact single-row and total MSEP wrappers, and has nonzero calendar CDR on three distinct updates: rows 0, 1, and 2 reveal their independent first-step shocks at calendar times 0, 1, and 2.*

Proof. Each row shock is Rademacher. The first residual has the prescribed second moment and every later residual vanishes. $\square$

Chapter 24

Röhr: arbitrary-horizon approximations

The primary-author presentation *Chain Ladder and Error Propagation*, given by Ancus Röhr at the CAE Fall 2014 meeting [11], states two relative-MSEP approximations on slide 13. If $\hat{u}_j^2$ is the squared relative uncertainty of factor j and $\hat{q}_j$ is its influence on the predicted ultimate $\hat{X}$, they are

$$\frac{\text{MSEP}}{\hat{X}^2} \approx \sum_{j \in S} \hat{u}_j^2 \hat{q}_j$$

for ultimate development and

$$\frac{\text{MSEP}(\text{CDR}_k)}{\hat{X}^2} \approx \sum_{j \in S} \hat{u}_j^2 \frac{\hat{q}_j - \hat{q}_{j-k}}{1 - \hat{q}_{j-k}}$$

for the claims development result over the next k years. The source sets $\hat{q}_j = 0$ below the active index range. Slide 17 labels an exact algebraic decomposition of each approximation as process and parameter error, and slide 20 shows the $k = 1$ form.

The formulas below are definitions because the source marks both MSEP statements with $\approx$. The Lean results prove algebra about those definitions. They do not construct $\hat{q}_j$ from its logarithmic derivative or identify an approximation with a stochastic MSEP theorem. The formulas are attributed to the 2014 slides, not to an inaccessible display of the 2016 paper.

24.1 Source formulas

Definition 24.1 (Slide 13, ultimate approximation). For a finite index set S , define the ultimate relative-MSEP approximation by $\sum_{j \in S} \hat{u}_j^2 \hat{q}_j$. Its slide-17 process and parameter components are $\sum_j \hat{u}_j^2 (1 - \hat{q}_j) \hat{q}_j$ and $\sum_j \hat{u}_j^2 \hat{q}_j^2$.

Definition 24.2 (Slide 13, arbitrary-horizon approximation). For current and earlier influences q and q^- , put

$$w(q, q^-) = \frac{q - q^-}{1 - q^-}.$$

The horizon approximation is $\sum_j \hat{u}_j^2 w(\hat{q}_j, \hat{q}_{j-k})$. Its displayed process and parameter terms are

$$\hat{u}_j^2 \frac{(1 - \hat{q}_j)(\hat{q}_j - \hat{q}_{j-k})}{(1 - \hat{q}_{j-k})^2} \quad \text{and} \quad \hat{u}_j^2 \left(\frac{\hat{q}_j - \hat{q}_{j-k}}{1 - \hat{q}_{j-k}} \right)^2.$$

Definition 24.3 (*k*-year and one-year forms). The *k*-year definition substitutes $q_j^- = q_{j-k}$ when $k \leq j$ and sets $q_j^- = 0$ when $j < k$. The one-year definition uses q_{j-1} when $1 \leq j$ and zero at the earlier index.

24.2 Exact splits

Theorem 24.4 (Slide 17, ultimate split). *The ultimate approximation splits exactly:*

$$\sum_j \hat{u}_j^2 \hat{q}_j = \sum_j \hat{u}_j^2 (1 - \hat{q}_j) \hat{q}_j + \sum_j \hat{u}_j^2 \hat{q}_j^2.$$

Proof. Expand the right side and collect terms. □

Theorem 24.5 (Slide 17, arbitrary-horizon split). *If $q_j^- \neq 1$ on S , then*

$$\sum_j \hat{u}_j^2 \frac{q_j - q_j^-}{1 - q_j^-} = \sum_j \hat{u}_j^2 \frac{(1 - q_j)(q_j - q_j^-)}{(1 - q_j^-)^2} + \sum_j \hat{u}_j^2 \left(\frac{q_j - q_j^-}{1 - q_j^-} \right)^2.$$

Proof. Put the two right-hand numerators over the common denominator $(1 - q_j^-)^2$ and sum the termwise identity. □

24.3 Endpoints and bounds

Theorem 24.6 (Horizon specializations). *Zero earlier influence reduces the horizon approximation and both split components to their ultimate forms. The $k = 1$ definition is the one-year formula. If $q_j \neq 1$ on S , the algebraic $k = 0$ endpoint is zero. More generally, if $q_{j-k} = 0$ for every active $j \in S$, the k -year formula is the ultimate formula.*

Proof. Substitute the indicated earlier influence in each definition. □

Theorem 24.7 (Influence bounds). *Suppose $0 \leq q_j^- \leq q_j \leq 1$ and $q_j^- < 1$. Then each horizon weight lies in $[0, 1]$. With $\hat{u}_j^2 \geq 0$, every horizon, process, and parameter term is nonnegative. The horizon approximation is at most the ultimate approximation when $q_j^- \geq 0$, $q_j \leq 1$, and $q_j^- < 1$ pointwise.*

Proof. The denominator $1 - q_j^-$ is positive. Division preserves the pointwise inequalities, squares are nonnegative, and finite sums preserve them. □

Chapter 25

R ChainLadder option semantics

This chapter pins the parameter-risk branch in R package `ChainLadder` 0.2.22 at source commit [41f4e949](#). Lean formalizes the arithmetic in lines 229–299; it does not execute R.

Definition 25.1 (MSE option and parameter-risk step). For prior squared parameter risk p , projected claim c , fitted factor f and $v = f.se_k^2$, the two methods use

$$c^2v + pf^2 \quad \text{for Mack}, \quad c^2v + pf^2 + pv \quad \text{for Independence.}$$

Equivalently, the multiplier of p is f^2 or $f^2 + v$.

Theorem 25.2 (The option changes one term). *At a fixed step, Independence minus Mack is exactly pv .*

Proof. Unfold the two constructors and subtract. □

Definition 25.3 (Per-row source recursion). Starting from zero on the latest diagonal, use $c = \hat{C}_{i,k}$, $f = \hat{f}_k$ and $v = \hat{\sigma}_k^2/S_k$.

Theorem 25.4 (Per-row option identities). *For nonzero fitted factors and $a_k = \hat{\sigma}_k^2/(\hat{f}_k^2 S_k)$,*

$$\begin{aligned} \text{Mack: } \hat{C}_{i,n-1}^2 \sum_k a_k &= \text{mackEstimation}, \\ \text{Independence: } \hat{C}_{i,n-1}^2 \left(\prod_k (1 + a_k) - 1 \right) &= \text{bbmwEstimation}. \end{aligned}$$

Proof. Induct along the row and use $\hat{\sigma}_k^2/S_k = \hat{f}_k^2 a_k$ and $\hat{C}_{i,k+1} = \hat{C}_{i,k} \hat{f}_k$. □

Definition 25.5 (Total source recursion). For total parameter risk, the package replaces c by the projected sum M_k .

Theorem 25.6 (Total recursion branches and closed form). *Mack propagates each increment through later f_l^2 ; Independence uses $f_l^2 + v_l$. The theorem gives the finite-sum closed form for both.*

Proof. Induct and split the last factor from each product. □

Definition 25.7 (Named row and total formulas). The constructors select the existing Mack or BMW row and total formulas.

Theorem 25.8 (Exact aggregate option difference). *The selected aggregate formulas differ by the existing row-wise BMW-minus-Mack remainder weighted by the projected ultimates.*

Proof. This is `bbmwTotalEstimation_sub_mackTotalEstimation`. □

Chapter 26

Stochastic Bühlmann-Straub credibility

Bühlmann and Straub [8], Sections 3, 5, and 6, model a loss ratio X_l using a latent risk parameter Θ . This chapter fixes the contract conditions. Conditional on Θ , the loss ratios have common mean $M = \mu(\Theta)$ and process variance $\sigma^2(\Theta)/P_l$, where P_l is exposure. The structural process parameter is $v = E[\sigma^2(\Theta)]$. Their Section 6 minimizes expected squared prediction error. This chapter proves the one-risk moment calculation and obtains the credibility factor in display (4), page 123. The scanned primary source is available from E-Periodica under DOI 10.5169/seals-967024.

Definition 26.1 (Conditional moment model). Let $\mathcal{G} = \sigma(\Theta)$, $S = \sigma^2(\Theta)$, and $e_l = X_l - M$. For the finite exposure sample, the model records

$$E[X_l | \mathcal{G}] = M, \quad E[e_l^2 | \mathcal{G}] = \frac{S}{P_l}, \quad E[e_l e_r | \mathcal{G}] = 0 \quad (l \neq r),$$

together with $E[S] = v$, $E[M] = m$, $E[(M - m)^2] = w$, positive exposures, measurability, and the needed integrability. The source assumes conditional independence; the Lean proof records only the conditional cross moments it consumes.

Definition 26.2 (Predictor, expected loss, and coefficient loss). Write $P = \sum_l P_l$, $\bar{X} = \sum_l P_l X_l / P$, and

$$\widehat{M}_z = m + z(\bar{X} - m), \quad R(z) = E[(M - \widehat{M}_z)^2].$$

The exposure-scaled coefficient loss is

$$Q(z) = Pw(1 - z)^2 + vz^2.$$

This stochastic objective is separate from the deterministic penalized sample loss in Definition 18.1.

Theorem 26.3 (Conditional residual moments). *The weighted residual total has conditional mean zero and conditional second moment SP . When the total exposure P is nonzero, it follows that*

$$E[(\bar{X} - M)^2] = \frac{v}{P}, \quad E[(M - m)(\bar{X} - M)] = 0.$$

Every predictor $m + z(\bar{X} - m)$ in the scalar family has expectation m , so it satisfies the source's expectation-unbiasedness constraint.

Proof. Scale each residual by its exposure. Conditional cross moments remove the off-diagonal terms, and $P_l^2(S/P_l) = SP_l$. Pulling the centered latent mean through conditional expectation removes the remaining cross term. $\square$

Theorem 26.4 (Expected prediction loss). *For every scalar coefficient z and nonzero total exposure,*

$$R(z) = (1 - z)^2 w + z^2 \frac{v}{P}, \quad PR(z) = Q(z).$$

Thus the deterministic coefficient quadratic is derived from the expected squared prediction loss, rather than postulated.

Proof. Expand $M - \widehat{M}_z = (1 - z)(M - m) - z(\bar{X} - M)$ and apply Theorem 26.3. $\square$

Theorem 26.5 (Credibility is the risk minimizer). *If $v + Pw > 0$, then*

$$Q(z) = Q(Z) + (v + Pw)(z - Z)^2, \quad Z = \frac{Pw}{v + Pw}.$$

For $P > 0$, Z minimizes the actual expected squared prediction loss. The predictor $m + Z(\bar{X} - m)$ is pointwise equal to `buhlmannStraubEstimate`. Supplying the source's empirical portfolio mean as m gives its display (4) estimator.

Proof. Complete the square in Q , use positivity to compare scaled expected losses, and apply the deterministic credibility-form identity. $\square$

Theorem 26.6 (Finite process-shock witness). *There is a two-outcome probability model with one unit exposure, constant latent mean 1, and observed loss ratio 0 or 2 with equal probability. It satisfies the moment model with process variance $v = 1 > 0$ and has a nonconstant observation. Its between-risk variance is zero; the general theorems do not require that specialization.*

Proof. Conditional expectation given the trivial σ -algebra is the two-point average. $\square$

Chapter 27

The stochastic Bornhuetter-Ferguson model

Mack, *The Prediction Error of Bornhuetter/Ferguson* [9], states the increment model on page 91 and derives the prediction-error calculation on pages 95–99. The exact probability statements and the practical assessments are separated below.

Definition 27.1 (BF1–BF3 and the future reserve). For zero-based increments S_{ik} , BF1 says that the full family is independent. On a finite horizon m , BF2 and BF3 say

$$E[S_{ik}] = x_i y_k, \quad \text{Var}(S_{ik}) = x_i \sigma_k^2, \quad \sum_{k < m} y_k = 1.$$

The true reserve after the first d periods is $R_i = \sum_{d \leq k < m} S_{ik}$, and $z_d = \sum_{k < d} y_k$ is the expected reported proportion.

Theorem 27.2 (Mean and process variance of a BF reserve). *Under BF2 and integrability,*

$$E[R_i] = x_i \sum_{d \leq k < m} y_k.$$

If also $d \leq m$, this equals $x_i(1 - z_d)$. Under BF1, BF3, and square integrability,

$$\text{Var}(R_i) = x_i \sum_{d \leq k < m} \sigma_k^2.$$

Proof. The integral commutes with the finite sum. Pairwise independence supplied by BF1 makes the finite-sum variance additive. $\square$

Definition 27.3 (Random BF estimate and MSEP). For a random initial ultimate U and a random selected cumulative pattern z , the BF estimate is $\hat{R}^{BF} = U(1 - z)$. The unconditional MSEP is $E[(\hat{R}^{BF} - R)^2]$. Mack’s source conditions on the observed increments; common independence means that the pair consisting of estimate and future reserve is independent of the observed-data sigma-algebra.

Theorem 27.4 (Product variance and the single-row MSEP). *If U and z are independent and the displayed products are square-integrable, then*

$$\text{Var}(U(1 - z)) = (E[U]^2 + \text{Var}(U)) \text{Var}(z) + \text{Var}(U)(1 - E[z])^2.$$

For $i < n$ under BF1, BF3, the stated square-integrability assumptions, and independence and equal means of the BF estimate and true future reserve,

$$\text{mse}(\widehat{R}^{BF}) = \text{Var}(\widehat{R}^{BF}) + x_i \sum_{d \leq k < m} \sigma_k^2.$$

Under common independence from an observed-data sigma-algebra, square integrability, and measurability of the prediction-target pair, the conditional expectation of the squared prediction error equals that same constant almost surely.

Proof. Factor first and second product moments by independence. For MSEP, center the prediction error and add the variances of the independent estimate and target. $\square$

Definition 27.5 (Raw pattern and variance estimators). On a finite set A of accident years with known volumes x_i , equations (1) and (2) become

$$\widehat{y}_k = \frac{\sum_{i \in A} S_{ik}}{\sum_{i \in A} x_i}, \quad \widehat{\sigma}_k^2 = \frac{1}{|A| - 1} \sum_{i \in A} \frac{(S_{ik} - x_i \widehat{y}_k)^2}{x_i}.$$

A general linear estimator is $\sum_{i \in A} a_i S_{ik}$.

Theorem 27.6 (Mack equations (1) and (2)). *For nonnegative volumes with nonzero total, the raw estimator is unbiased and has variance $\sigma_k^2 / \sum_i x_i$. If $\sigma_k^2 \geq 0$, every linear unbiased competitor, $\sum_i a_i x_i = 1$, has at least that variance. With at least two contributors and nonzero individual volumes, the residual estimator is unbiased for σ_k^2 .*

Proof. Independence makes the linear-estimator variance a diagonal quadratic form. Complete its square around $a_i = 1 / \sum_j x_j$. For the residual estimator, subtract the fitted-mean square from the true-centre sum; their expectations are σ_k^2 times $|A|$ and one, respectively. $\square$

Definition 27.7 (Practical assessments). Equations (3), (4), (6), and (7) define the raw development proportions, raw column variances, squared standard errors of the selected proportions, and the minimum of the prefix and suffix variance sums. The single-year plug-in is

$$\widehat{\text{mse}} = \widehat{\text{Var}}(R_i) + (U_i^2 + s_U^2) s_z^2 + s_U^2 (1 - z)^2.$$

The total-reserve covariance assessment on page 99 is also a definition. Its correlations, standard errors, smoothing, and tail values are supplied by the actuary. No unbiasedness or consistency claim is attached to those selected or assessed inputs.

The standalone equation (3), (4), and (6) definitions retain Mack's one-based symbols $\mathbf{n}$ and $\mathbf{k}$; the stochastic theorems above instead use an explicit finite contributor set and therefore follow the repository's zero-based triangle convention.

Theorem 27.8 (Assessment algebra and total covariance). *The single-year assessment is exactly its named process term plus its two estimation terms. Before Mack omits the middle term for the total-reserve approximation, independence of the pairs $\{X, W\}$ and $\{Y, Z\}$ gives*

$$\text{Cov}(XY, WZ) = \text{Cov}(X, W) E[Y] E[Z] + \text{Cov}(X, W) \text{Cov}(Y, Z) + E[X] E[W] \text{Cov}(Y, Z).$$

Proof. Factor the required cross moments between the two independent pairs, substitute the covariance identities, and normalize the polynomial expression. $\square$

Theorem 27.9 (A finite BF witness). *On the one-point probability space, take $C_{i,k} = k + 1$, so every increment is one, $x_i = y_k = 1$, and $\sigma_k^2 = 0$ on the one-period horizon. BF1–BF3 and the normalization hold together. The future reserve has mean one and variance zero, and prediction by one has conditional MSEP zero. This checks satisfiability of the hypotheses; it does not provide a positive-variance BF model.*

Proof. Every random variable is constant. The generated sigma-algebras are trivial, so the full increment family is independent. $\square$

Chapter 28

Panjer recursion

Panjer, *Recursive Evaluation of a Family of Compound Distributions* [10], equation (1), assumes claim-count probabilities satisfying $p_n = p_{n-1}(a + b/n)$. For a severity distribution supported on the positive integers, equation (5) gives a direct recursion for the aggregate mass.

Definition 28.1 (Formal generating series and compound mass). For a coefficient sequence p , let $P(X) = \sum_n p_n X^n$. For a severity sequence f with $f_0 = 0$, the compound series is $G(X) = P(F(X))$. The predicate **PanjerFrequency** is the division-free recurrence

$$(n+1)p_{n+1} = (a(n+1) + b)p_n.$$

The definition **panjerEquation5** is the finite sum printed in the source.

Theorem 28.2 (Panjer equation (5)). *The count recurrence is equivalent to*

$$P' = (a+b)P + aXP'.$$

Formal substitution and the power-series chain rule give, for $i \geq 1$,

$$g_i = \sum_{j=1}^i \left(a + b \frac{j}{i} \right) f_j g_{i-j}, \quad g_0 = p_0.$$

The coefficient of $P(F(X))$ is also the convolution mixture in the source's equation (6).

Proof. Convert the count recurrence to a formal differential equation, substitute the severity series, apply the formal chain rule, and compare coefficients. $\square$

Theorem 28.3 (Geometric and unit-severity check). *The geometric coefficients $(1-q)q^n$ satisfy the Panjer recurrence with $(a,b) = (q,0)$. For deterministic severity one, composition leaves the count mass unchanged, and the resulting geometric instance satisfies equation (5) at every index.*

Proof. Expand the geometric recurrence. The unit-severity series is X , so substitution by it is the identity. $\square$

Chapter 29

Non-vacuity: a nondegenerate model

Theorem 29.1 (A nondegenerate Mack model exists). *There is a probability space, a random triangle with three accident years, and constants f_k, σ_k^2 with $\sigma_0^2 > 0$, satisfying (M1), (M3) and (M2'), in which the first development step is genuinely random. The model: eight equally likely outcomes, an independent ± 1 shock per accident year, $C_{i,0} = c$ and $C_{i,k} = cf^k(1 + s\xi_i)$ for $k \geq 1$, with $\mathcal{D}_0$ trivial and $\mathcal{D}_k$ everything for $k \geq 1$.*

Proof. Conditional expectations given the trivial σ -algebra are eight-term averages; the shocks sum to zero, have unit square, and are orthogonal. $\square$

Theorem 29.2 (Headline theorems instantiated). *On that model, $E[\hat{f}_0 | \mathcal{D}_0] = 2$ although $\hat{f}_0$ is random; $E[\hat{C}_{2,2} | \mathcal{D}_0] = E[C_{2,2} | \mathcal{D}_0]$; $E[(\hat{f}_0 - 2)^2 | \mathcal{D}_0] = \sigma_0^2/S_0$; and $E[\hat{\sigma}_0^2 | \mathcal{D}_0] = \sigma_0^2 = 4$.*

Proof. Direct instantiation of the general theorems; every hypothesis is discharged. $\square$

Chapter 30

Reference computation

Mack’s 1993 example [1], based on the Taylor and Ashe triangle [19], is reproduced to every printed digit (nine reserves, the total 18,681 thousand, all standard-error percentages) by a dependency-free script implementing the definitions above in floating point. The Lean definitions are over $\mathbb{R}$ and noncomputable; the script is the reference computation, not the Lean code.

Bibliography

- [1] T. Mack, Distribution-free calculation of the standard error of chain ladder reserve estimates, *ASTIN Bulletin* 23 (1993) 213-225.
- [2] T. Mack, The standard error of chain ladder reserve estimates: recursive calculation and inclusion of a tail factor, *ASTIN Bulletin* 29 (1999) 361-366.
- [3] G. Quarg and T. Mack, Munich chain ladder: a reserving method that reduces the gap between IBNR projections based on paid losses and IBNR projections based on incurred losses, *Blätter der DGVM* 26 (2004) 597-630; authorized English reprint in *Variance* 2 (2008) 266-299.
- [4] P. England and R. Verrall, Analytic and bootstrap estimates of prediction errors in claims reserving, *Insurance: Mathematics and Economics* 25 (1999) 281-293.
- [5] P. England and R. Verrall, Stochastic claims reserving in general insurance, *British Actuarial Journal* 8 (2002) 443-518.
- [6] P. England, Addendum to “Analytic and bootstrap estimates of prediction errors in claims reserving”, *Insurance: Mathematics and Economics* 31 (2002) 461-466.
- [7] M. Merz and M. V. Wüthrich, Modelling the claims development result for solvency purposes, *Casualty Actuarial Society E-Forum*, Fall 2008, 542-568.
- [8] H. Bühlmann and E. Straub, Glaubwürdigkeit für Schadensätze, *Mitteilungen der Vereinigung Schweizerischer Versicherungsmathematiker* 70 (1970) 111-133.
- [9] T. Mack, The prediction error of Bornhuetter/Ferguson, *ASTIN Bulletin* 38 (2008) 87-103.
- [10] H. H. Panjer, Recursive evaluation of a family of compound distributions, *ASTIN Bulletin* 12 (1981) 22-26.
- [11] A. Röhr, Chain ladder and error propagation, presentation at the Casualty Actuaries in Europe Fall Meeting, 31 October 2014.
- [12] A. E. Renshaw and R. J. Verrall, A stochastic model underlying the chain-ladder technique, *British Actuarial Journal* 4 (1998) 903-923.
- [13] T. Mack, A simple parametric model for rating automobile insurance or estimating IBNR claims reserves, *ASTIN Bulletin* 21 (1991) 93-109.
- [14] M. V. Wüthrich and M. Merz, *Stochastic Claims Reserving Methods in Insurance*, Wiley, Chichester, 2008.

- [15] M. Buchwalder, H. Bühlmann, M. Merz and M. V. Wüthrich, The mean square error of prediction in the chain ladder reserving method (Mack and Murphy revisited), *ASTIN Bulletin* 36 (2006) 521-542.
- [16] D. M. Murphy, Unbiased loss development factors, *Proceedings of the Casualty Actuarial Society* 81 (1994) 154-222.
- [17] T. Mack, Measuring the variability of chain ladder reserve estimates, *Casualty Actuarial Society Forum*, Spring 1994, 101-182.
- [18] A. Röhr, Chain ladder and error propagation, *ASTIN Bulletin* 46 (2016) 293-330.
- [19] G. C. Taylor and F. R. Ashe, Second moments of estimates of outstanding claims, *Journal of Econometrics* 23 (1983) 37-61.